

JC2001 - Introduction to Software Engineering

Group Project · Technical Report

Academic year 2026-27

Campus Market: Design, Implementation and Evaluation of an On-Campus Second-Hand Trading Platform

*A proof-of-concept web application for buying, selling and swapping second-hand goods
within a single university campus*

Group number: 9

Programme of study: BSc (Artificial Intelligence)

Group member	Name	Student ID
Group member 1	<Full name>	<Student ID>
Group member 2	<Full name>	<Student ID>
Group member 3	<Full name>	<Student ID>
Group member 4	<Full name>	<Student ID>
Group member 5	<Full name>	<Student ID>

Submitted to the MyAberdeen course assessment area · Submission date: <date>

Table of Contents

Table of Contents	2
List of Figures	5
List of Tables	7
1 Introduction and Project Context	9
1.1 Background	9
1.2 The Problem	9
1.3 Project Objectives	10
1.4 Intended Users and Stakeholders	10
1.5 Scope of the Proof of Concept	11
1.6 Main Features	11
1.7 Potential Practical and Business Value	12
1.8 Structure of This Report	12
2 Project Approach and Method	13
2.1 Development Process	13
2.2 Requirements Elicitation	13
2.3 Tools and Configuration Management	14
2.4 Team Organisation	14
2.5 Constraints and Assumptions	15
2.6 How the Work Was Verified	15
3 Requirements and Use Cases	17
3.1 Requirements Modelling Approach	17
3.2 Functional Requirements	17
3.3 Non-Functional Requirements	18
3.4 Requirement Prioritisation	19
3.5 System Context	19
3.6 Use Case Model	20
3.7 Detailed Use Case Descriptions	22
3.8 Business Process Model	25
3.9 Activity Model for Publishing	26

3.10 Requirements Traceability	27
4 System Design and Architecture.....	28
4.1 Architectural Style and the Alternatives Considered.....	28
4.2 Layered Architecture	28
4.3 Module Decomposition.....	29
4.4 Domain Model	30
4.5 Database Design.....	31
4.6 Transaction State Machine.....	33
4.7 Interaction Design.....	34
4.8 Interface Design	36
4.9 Client-Side Design	38
4.10 Design Decisions and Trade-offs.....	39
5 Implementation	41
5.1 Technology Stack.....	41
5.2 Software Structure and Size.....	41
5.3 Accounts, Sessions and Credential Storage	42
5.4 The Catalogue Query	42
5.5 Photograph Upload and Storage	43
5.6 The Status Machine.....	43
5.7 The Conversation List Query	44
5.8 Client-Side Implementation	45
5.9 Design System	45
5.10 User Interface.....	46
5.11 Coding Standards and Conventions	53
5.12 Deployment and Reproducibility	54
6 Testing and Evaluation	55
6.1 Test Strategy	55
6.2 Test Environment and Harness	55
6.3 Unit and Component Test Results	56
6.4 Interface Testing	57
6.5 Negative and Boundary Testing.....	57

6.6 State Machine Testing.....	57
6.7 Authorisation Testing.....	58
6.8 Requirements Coverage	58
6.9 Performance Testing	59
6.10 Scaling Measurement.....	61
6.11 Interface and Usability Evaluation.....	62
6.12 Verification of the Non-Functional Requirements	62
6.13 Defects Found and Fixed	62
6.14 Threats to Validity and What Was Not Tested	62
7 Conclusions and Future Work	64
7.1 Outcomes Against the Objectives	64
7.2 Limitations	64
7.3 Future Work	65
7.4 Reflection on the Process.....	66
7.5 Closing Statement	66
References.....	68
Appendix A - Group Workload Profile	70
A.1 Quantitative contribution	70
A.2 Qualitative contribution	70
A.3 Declaration	70
Appendix B - Requirements Traceability Matrix	72
Appendix C - Verification Detail.....	74
C.1 Interface contract (Table 6.3)	74
C.2 State machine transition coverage (Table 6.5)	74
C.3 Authorisation matrix (Table 6.6).....	75
C.4 Verification of the non-functional requirements (Table 6.9)	75
C.5 Defect log (Table 6.10)	76

List of Figures

Figure 3.1	System context: the actors, the browser and the local filesystem the proof of concept interacts with	20
Figure 3.2	Use case diagram: eighteen use cases across two actors, with the visitor actor generalising to the registered student	21
Figure 3.3	Business process model for a complete second-hand transaction, with the server-side guards shown in the system lane	26
Figure 3.4	Activity model for publishing a listing, showing the three decisions and the duplicated client-side and server-side enforcement	27
Figure 4.1	Layered architecture, showing the endpoint groups of the application layer and the components of the data layer	29
Figure 4.2	Module decomposition and the direction of each dependency	30
Figure 4.3	Domain class diagram, showing the six entities, their attributes and their associations	31
Figure 4.4	Entity-relationship diagram of the six-table schema, with foreign keys and cardinality	32
Figure 4.5	State machine governing the life cycle of a listing, with the actor and the guard for each transition	33
Figure 4.6	Sequence diagram for publishing a listing, from form submission to committed rows	34
Figure 4.7	Sequence diagram for reserving an item and confirming the sale, including the two guards and the rejection path	35
Figure 4.8	Sequence diagram for the conversation list, opening a thread and sending a reply	36
Figure 4.9	The JSON endpoint surface, grouped by resource, with the methods and the session requirement of each endpoint	37
Figure 4.10	The client-side view router: five views, the persistent header, and the guards on the views that require a session	39
Figure 5.1	Technology choices and the reasoning behind each one	41
Figure 5.2	How a status request is validated before the row is written, driven by a five-entry rule table	44
Figure 5.3	Rendering the conversation list with one aggregate query instead of a query per conversation	44
Figure 5.4	The catalogue grid: status badge, cover image, price, seller and category on each card	47
Figure 5.5	The empty search result: the condition is named and a next step offered, instead of leaving the grid blank	47
Figure 5.6	The detail view as a signed-in buyer sees it: reserve, save and message, with the seller and the structured attributes	48
Figure 5.7	The same view for the seller: the buying actions are replaced by the management actions	49
Figure 5.8	The shared confirmation dialog used for every destructive action in the interface	49
Figure 5.9	The publishing form: chip selectors for category and condition, a photograph drop zone, and validation messages attached to the field that failed	50
Figure 5.10	The personal centre: profile summary with counts, and the student's own listings	51
Figure 5.11	An open conversation: own and other messages distinguished, relative times, and the composer anchored to the bottom	52

Figure 5.12	The interface at a 400-pixel viewport: the grid collapses to two columns and the header reduces to its essential controls	53
Figure 6.1	Automated test cases by component. All 80 cases pass.	56
Figure 6.2	Requirements realised and independently verified, by category	59
Figure 6.3	In-process response time per operation at 500 listings, median and 95th percentile	60
Figure 6.4	Median response time for three catalogue read paths as the number of listings grows from 100 to 2,000	61

List of Tables

Table 1.1	Stakeholders and what each one needs from the system	10
Table 1.2	Scope boundary of the proof of concept	11
Table 1.3	Main features of the proof of concept and where each is specified and verified	11
Table 2.1	Roles and responsibilities. Individual contributions and the agreed percentage weightings are recorded in the Group Workload Profile in Appendix A.	14
Table 2.2	Project constraints and assumptions	15
Table 3.1	Functional requirements	17
Table 3.2	Non-functional requirements and how each is verified	18
Table 3.3	Prioritisation of requirements	19
Table 3.4	Use case summary	21
Table 4.1	Architectural alternatives and the reason each was rejected	28
Table 4.2	Responsibility of each layer, and what each layer is forbidden to do	29
Table 4.3	Module inventory	30
Table 4.4	Table specifications	32
Table 4.5	Status transition rules. The role and the source state are both required; a request that fails either check changes nothing.	33
Table 4.6	Endpoint catalogue. Session means that the endpoint refuses an unauthenticated request with 401 without evaluating the body.	37
Table 4.7	The response envelope shared by every endpoint	38
Table 4.8	Design decision log	39
Table 5.1	Component versions and roles	41
Table 5.2	Size of each component of the delivered software	42
Table 5.3	Validation rules applied to an upload or a listing, and where each is enforced	43
Table 5.4	Client-side rendering paths and the escaping applied to each	45
Table 5.5	Design token groups and what each controls	45
Table 5.6	Coding conventions applied across the codebase	53
Table 6.1	Test levels, techniques and detection capability	55
Table 6.2	Automated test suite by group	56
Table 6.4	Negative and boundary cases written deliberately	57
Table 6.7	In-process response time per operation. Dataset: 500 listings, 22 accounts, 61 conversations, one 40-message thread. Python 3.13.14, 200 repetitions for reads and 100 for writes.	60
Table 6.8	Median response time for three catalogue read paths as the dataset grows	61
Table 7.1	Outcome against each project objective	64
Table 7.2	Limitations of the delivered proof of concept	64
Table 7.3	Prioritised future work	65

1 Introduction and Project Context

1.1 Background

Every academic year a substantial volume of usable goods moves between students. Textbooks bought for a single semester, laptops replaced on an upgrade, monitors, bicycles, calculators, sports kit and kitchen equipment all become surplus to their owners while demand for those exact items exists a few hundred metres away in the next year group. Research on consumer resale consistently finds that the barriers to second-hand trade are not price but friction: the effort of listing, the uncertainty of whether anyone wants the item, and the difficulty of meeting a buyer (Kleppmann, 2017, discusses the same friction in the context of marketplaces more generally). On a university campus the physical distance is negligible, which makes the remaining friction purely informational. That is the gap this project addresses.

1.2 The Problem

Three channels currently carry almost all campus trade, and each fails in a different way.

- **General-purpose marketplaces.** These are designed for buyers and sellers who are strangers to one another and geographically dispersed. Listings are drowned in results from the wider city; handover must be negotiated with a counterparty who may be an hour away; and the platform charges a commission or sells promotion. For a student selling a 40-yuan textbook the net benefit is close to zero, so the channel is not used.
- **Social media groups and group chats.** This is where most transactions actually happen, but a chat stream is a poor catalogue. A listing is pushed out of view within hours, there is no search, no filtering, no way to see whether an item is still available, and no record that anything was ever sold. The same item is re-posted repeatedly by different people because there is no canonical listing.
- **Noticeboards and word of mouth.** These reach only the people who happen to walk past, and give a buyer no way to compare prices or assess condition before making a journey.

Two of the failures are specific to this setting and explain why a general-purpose platform cannot simply be pointed at a campus. The first is **trust**. Students transact with people they may share a lecture theatre with but do not know, and a platform open to the world gives them no signal that the counterparty is a member of the same institution and no agreed, public place to meet. The second, and the one this project treats as the central design problem, is **state**. In an informal channel there is no authoritative answer to the question "is this still

available?" Sellers answer that question dozens of times, buyers travel to inspect items that were sold yesterday, and neither party can tell whether an item is on sale, held for someone, or gone.

The economic and environmental consequences are straightforward. Sellers abandon items they cannot shift or give them away, recovering none of the value; buyers pay full price for new goods because they never saw the used alternative; and the institution has to dispose of a volume of usable goods that its own students would have bought. Addressing the problem matters because the affected population is large, price-sensitive, and already motivated to participate.

1.3 Project Objectives

The project set out to achieve five objectives, each stated so that its completion is demonstrable rather than asserted.

- (1) To specify the requirements of a campus second-hand trading platform, by identifying its users and stakeholders and producing a prioritised set of functional and non-functional requirements against which the remainder of the project can be traced.
- (2) To design a system architecture and data model for the platform, expressed as a layered architecture, a domain model, a relational schema and a set of interaction models, with the key design decisions and their alternatives recorded.
- (3) To implement a working proof of concept that demonstrates the core features end to end.
- (4) To verify the proof of concept against the stated requirements through automated unit, component and interface testing, negative and boundary testing, performance measurement and an interface evaluation, reporting the results honestly including cases that fail or remain unverified.
- (5) To deliver the documentation and presentation required by the module, consistent with the delivered software.

Chapter 7 returns to these objectives and states, for each, what was achieved and what was not.

1.4 Intended Users and Stakeholders

Table 1.1 *Stakeholders and what each one needs from the system*

Stakeholder	Relationship to the system	Primary need
Visitor (unregistered student)	Direct user. Reaches the site from a link or a search and has no account.	To browse and search the catalogue, read a listing in full, and see enough of the seller to decide whether to register.
Registered student	Direct user, acting in two roles. The same account sells items and buys them; the roles are separated only for the duration of a transaction.	As a seller: to publish quickly, to be reachable, and to stop receiving enquiries once an item has gone. As a buyer: to find what is available, to know whether it is

Stakeholder	Relationship to the system	Primary need
		still available, and to contact the seller.
The institution	Indirect stakeholder. Provides the campus, the students and, in a production setting, the hosting.	That the platform carries no money, holds the minimum personal data, and imposes no operational cost or regulatory obligation.
Course assessor	Indirect stakeholder. Installs, runs and evaluates the software.	That the software installs and runs on a clean machine without specialist knowledge, and that the implementation matches what the report claims.

1.5 Scope of the Proof of Concept

Table 1.2 *Scope boundary of the proof of concept*

In scope	Deliberately out of scope
Account registration, authentication and session handling	Institutional single sign-on and identity verification against a student record system
Listing creation with up to three photographs and a structured attribute set	Rich media (video), bulk import, and catalogue synchronisation from other platforms
Catalogue browsing with keyword search, category and status filters, three sort orders and pagination	Recommendation, personalisation and relevance ranking
Listing detail with gallery, description, seller identity and a public comment thread	Reputation scores, reviews and dispute history
Favourites, private messaging and a personal centre	Push or email notification, and real-time delivery
Server-validated transaction status: reserve, confirm, cancel, take offline, relist	Online payment, escrow, delivery logistics and refunds
A single-node deployment on a local machine	Horizontal scaling, high availability and multi-tenant operation

1.6 Main Features

Each feature below maps onto a group of functional requirements specified in Chapter 3 and onto the test cases that verify it in Chapter 6.

Table 1.3 *Main features of the proof of concept and where each is specified and verified*

Feature	What it does	Requirements	Verified by
Accounts and sessions	Register, log in, log out, retrieve the current profile; passwords stored as a salted hash	FR-01 to FR-04	TC-A01 to TC-A10
Publishing	Create, edit and soft-delete a listing, with up to three photographs and validation of all required fields	FR-05 to FR-10	TC-L01 to TC-L12, TC-U01 to TC-U08
Catalogue	Paginated listing grid with keyword search across two columns, category and status filters and three sort orders	FR-11 to FR-15	TC-C01 to TC-C12
Listing detail	Photo gallery, description, seller identity, structured attributes, comment thread and role-aware actions	FR-16 to FR-18	TC-C13, TC-C14, TC-S05 to TC-S08
Favourites	Save and unsave a listing; a saved list in the personal centre	FR-19 to FR-21	TC-S01 to TC-S04
Transaction state	Reserve, confirm the sale, cancel a reservation, take offline and relist, each guarded by role and current state	FR-22 to FR-24	TC-T01 to TC-T15
Private messaging	Send a message, list conversations grouped by	FR-25 to FR-27	TC-S09 to TC-S16

Feature	What it does	Requirements	Verified by
	partner, open a thread and clear unread state		
Personal centre	My listings, my favourites, my sold items and my conversations in one place	FR-28	TC-P01 to TC-P05

1.7 Potential Practical and Business Value

The value of the system comes from removing information friction rather than from introducing a new capability. Students already want to trade; what they lack is a place where a listing persists, is searchable, and answers the availability question authoritatively.

For students, the practical value is measurable in money and time. A textbook resold at half price saves the buyer half the purchase price and returns half to the seller; across a cohort of several hundred students the aggregate effect is significant. Time is saved at both ends: the buyer searches rather than walks a corridor of noticeboards, and the seller answers "still available?" once, in public, instead of once per enquirer. For the institution, a campus marketplace reduces the volume of usable goods sent to disposal at the end of the academic year and provides a student-life service that costs nothing to operate, because the software carries no payment and therefore no commission, no financial regulation and no personal financial data.

1.8 Structure of This Report

Chapter 2 describes the process used to produce the work and the constraints it ran under. Chapter 3 specifies the functional and non-functional requirements, the use case model and the process and activity models. Chapter 4 presents the architecture, the domain and data models, the interaction models and the design decisions, each with the alternatives that were rejected. Chapter 5 describes the implementation, including the technologies, the software structure, the algorithms that carry the most engineering risk and the user interface. Chapter 6 reports the testing and evaluation, with the results of eighty automated test cases, a per-operation performance measurement and a scaling measurement, and states what the testing does not cover. Chapter 7 concludes against the objectives of section 1.3 and sets out the limitations and the future work. Appendix A contains the group workload profile and Appendix B the full requirements traceability matrix.

2 Project Approach and Method

2.1 Development Process

The project followed an **incremental, prototype-driven** process rather than a single waterfall pass or a full agile ceremony. The reasoning was specific to the situation. The team had a fixed end date, a fixed team and a scope it did not yet understand: the requirement that turned out to be hardest to get right - what happens when two buyers want the same item - is not visible from a requirements interview alone. Building a narrow end-to-end prototype first made that requirement concrete, and the specification in Chapter 3 was written afterwards with the prototype in hand (Sommerville, 2016, describes this as evolutionary prototyping; Ries, 2011, makes the same argument for a minimum viable product).

The process had four characteristics.

- **Vertical slices rather than layers.** Each increment delivered a user-visible capability through every layer - interface, endpoint, query, storage - rather than completing one layer at a time. This kept the software runnable at every point and meant the team never had to integrate several untested layers at the end.
- **A written specification stabilised after the first prototype.** Early increments informed the requirements; later increments were built against the specified requirement identifiers, so that every feature can be traced back to a requirement and forward to a test.
- **Automated tests written alongside the code.** The suite described in Chapter 6 grew with the implementation and was run after every change, which is what allowed a late system-wide change - rewriting the front end - to be made without a regression.
- **A single shared repository with frequent commits.** All work went through Git, so changes are attributable and every increment is recoverable.

2.2 Requirements Elicitation

Requirements came from five sources, chosen to give different kinds of information rather than to be exhaustive. The most useful turned out to be the least formal: watching where an existing second-hand transaction actually broke down.

- **Semi-structured interviews** with six students across two year groups, selected for having bought or sold something second-hand in the previous year. These produced the failure modes of the current channels and the finding that "is it still available?" is the dominant interaction cost.

- **Document and artefact analysis** of four active campus resale group chats, sampled over two weeks. This produced the evidence for the re-posting problem and for the category vocabulary the interface now offers as fixed choices.
- **Task observation** of three listings followed from posting to handover, which produced the sequence of steps that became the business process model in section 3.8.
- **Supervisor discussion** in two scheduled sessions, which produced the scope reduction: dropping payment, delivery and reputation, each of which would have made the project undeliverable.
- **Competitive review** of public listings on two general-purpose marketplaces, which produced the attribute set buyers screen on: price, condition and photographs.

2.3 Tools and Configuration Management

- **Git and GitHub** for version control, so every change is attributable and every increment recoverable.
- **Python 3.8 or later** as the implementation language, because it is already present on the lab and staff machines and needs no runtime licence.
- **Flask 2.3.3** as the web framework, because routing, signed sessions and JSON handling arrive in one small dependency with no configuration ceremony.
- **SQLite 3** as the database, because it is file-based and the schema therefore travels with the submission.
- **pytest** for automated testing, because its fixture model gives per-test database isolation without extra libraries.
- **Playwright driving Microsoft Edge** for interface verification, because it reuses a browser that is already installed rather than downloading one.
- **Microsoft Word** for document production, which is the required output format.

2.4 Team Organisation

Work was divided by **feature vertical** rather than by technical layer. Splitting by layer would have produced three people waiting on each other with nothing runnable until late in the project, whereas splitting by feature gave every member a slice they could complete and demonstrate alone. A rotating reviewer checked each change against the requirement it claimed to satisfy.

Table 2.1 *Roles and responsibilities. Individual contributions and the agreed percentage weightings are recorded in the Group Workload Profile in Appendix A.*

Role	Held by	Responsibility
Requirements lead	Group member 1	Elicitation, functional requirements, requirement priority and the traceability matrix
Architecture lead	Group member 2	Architecture selection, module boundaries, interface contract,

Role	Held by	Responsibility
		supervisor liaison
Data and security lead	Group member 3	Schema design and normalisation, password storage, upload validation, performance measurement
Process and state lead	Group member 4	Transaction state machine, authorisation rules, negative and boundary test design
Interface lead	Group member 5	Design system, all views and feedback states, responsive behaviour, user manual
Reviewer (rotating)	All members	Checking each change against the requirement it claims to satisfy before it is merged

2.5 Constraints and Assumptions

Table 2.2 *Project constraints and assumptions*

#	Constraint or assumption	Consequence for the design
C1	The proof of concept must run on the assessor's machine with a single command and no configuration	Rules out a server-based database, a container runtime and a front-end build step. Drives the choice of SQLite and framework-free client code.
C2	The platform must not handle money	Excludes payment, escrow and refunds, and with them financial regulation, dispute resolution and the storage of financial data.
C3	No institutional identity provider is available to the project	Authentication is self-contained: username, password and an optional student identifier that the platform records but cannot verify.
C4	The project runs for a single academic year with a part-time team	Scope is fixed at the transaction loop for one campus; recommendation, notification and reputation are explicitly deferred.
C5	One campus, therefore a small and bounded user population	A single-node database is adequate. Capacity planning for a multi-campus deployment is out of scope and is listed as future work.
A1	Assumed: students have a browser and a campus network connection	No native client is required; a responsive web client serves both desktop and mobile.
A2	Assumed: buyers and sellers can meet in person on campus	The system models handover as an offline event and does not attempt to schedule or track delivery.
A3	Assumed: the volume of data is small in absolute terms	The absence of indexes beyond the primary keys is tolerable at this scale, and the scaling behaviour in section 6.10 confirms it, but the assumption is revisited in the limitations.

2.6 How the Work Was Verified

Verification was treated as part of the process rather than as a phase at the end, through three mechanisms used continuously: an automated suite of eighty test cases exercising every JSON endpoint including the negative and boundary paths, run before every push; a scripted interface check that drives the application in a real browser, visits every screen in both an authenticated and an anonymous state, and fails on a console error or an unexpected 4xx or 5xx response; and a written traceability matrix, maintained from the moment the

requirements were numbered. The matrix is in Appendix B and the results are reported in Chapter 6.

3 Requirements and Use Cases

3.1 Requirements Modelling Approach

Requirements are recorded as numbered, atomic statements, because numbering is what makes traceability possible: a requirement that cannot be referred to by name cannot be shown to have been implemented or tested. Functional requirements carry the prefix **FR**, non-functional requirements **NFR**, use cases **UC** and test cases **TC**. Every requirement is written with a single verb, states what the system does rather than how, and is verifiable by a test that either passes or fails (Sommerville, 2016; Pressman and Maxim, 2020).

3.2 Functional Requirements

Table 3.1 specifies the twenty-eight functional requirements, grouped by the feature they belong to. The final column names the use cases of section 3.6 in which the requirement is exercised.

Table 3.1 *Functional requirements*

ID	Requirement	Group	Use case
FR-01	The system shall allow a visitor to register an account with a username, a password and optional nickname, email and student identifier.	Accounts	UC-01
FR-02	The system shall authenticate a registered student against a username and password and establish a session.	Accounts	UC-02
FR-03	The system shall allow an authenticated student to terminate their session.	Accounts	UC-02
FR-04	The system shall return the authenticated student's own profile, excluding any credential material.	Accounts	UC-12
FR-05	The system shall allow an authenticated student to publish a listing with a title, description, price, category and condition.	Publishing	UC-06
FR-06	The system shall allow up to three photographs to be attached to a listing, each stored under a system-generated name.	Publishing	UC-07
FR-07	The system shall reject a listing that has no title, no price, or a price that is not numeric.	Publishing	UC-06
FR-08	The system shall allow the owner of a listing to edit its title, description, price, category and condition.	Publishing	UC-08
FR-09	The system shall allow the owner of a listing to take it offline without deleting its history or its photographs.	Publishing	UC-09
FR-10	The system shall allow the owner of an offline listing to put it back on sale.	Publishing	UC-09
FR-11	The system shall present listings as a paginated catalogue with a fixed page size and a total result count.	Catalogue	UC-03
FR-12	The system shall allow listings to be searched by a keyword matched case-insensitively against title and description.	Catalogue	UC-04
FR-13	The system shall allow listings to be filtered by category.	Catalogue	UC-04
FR-14	The system shall allow listings to be filtered by status.	Catalogue	UC-04
FR-15	The system shall allow listings to be ordered by recency, ascending price or descending price.	Catalogue	UC-04
FR-16	The system shall present a listing detail page containing the photograph gallery, description, price, category, condition, posting time and the seller's	Detail	UC-05

ID	Requirement	Group	Use case
	display name.		
FR-17	The system shall present the public comments on a listing in chronological order.	Detail	UC-10
FR-18	The system shall allow an authenticated student to post a comment on a listing.	Detail	UC-10
FR-19	The system shall allow an authenticated student to save a listing as a favourite.	Favourites	UC-11
FR-20	The system shall allow an authenticated student to remove a favourite.	Favourites	UC-11
FR-21	The system shall list the authenticated student's favourites, most recently saved first.	Favourites	UC-11
FR-22	The system shall allow an authenticated student who is not the seller to reserve a listing that is on sale, moving it to reserved.	Transaction	UC-13
FR-23	The system shall allow the seller of a reserved listing to confirm the sale, moving it to sold.	Transaction	UC-14
FR-24	The system shall allow the seller of a reserved listing to cancel the reservation, returning it to on sale.	Transaction	UC-15
FR-25	The system shall allow an authenticated student to send a private message to another registered student.	Messaging	UC-16
FR-26	The system shall present the authenticated student's conversations grouped by the other participant, showing the most recent message of each.	Messaging	UC-17
FR-27	The system shall allow the authenticated student to open a conversation, read its messages in chronological order, and have received messages marked as read.	Messaging	UC-18
FR-28	The system shall present a personal centre containing the authenticated student's own listings, favourites, sold items and conversations.	Personal centre	UC-12

3.3 Non-Functional Requirements

Table 3.2 *Non-functional requirements and how each is verified*

ID	Attribute	Requirement	Verification
NFR-01	Usability	A student who has never used the system shall be able to publish a listing without written instructions or training.	Task observation with three first-time users (section 6.11)
NFR-02	Usability	Every action shall produce visible feedback in the interface, and no destructive action shall proceed without an explicit confirmation.	Interface inspection; scripted capture of every screen (section 6.11)
NFR-03	Usability	All views shall share one visual language: a single type scale, one spacing scale and one set of component styles.	Review of the design tokens against the component stylesheet
NFR-04	Accessibility	Every interactive control shall be reachable and operable by keyboard, with a visible focus indicator, and modal surfaces shall expose their role and title to assistive technology.	Keyboard-only walkthrough of each view; inspection of ARIA attributes
NFR-05	Accessibility	The interface shall remain usable without horizontal scrolling from a 430-pixel mobile viewport to a 1440-pixel desktop viewport.	Capture of both viewport sizes (figures 5.15 and 5.16)
NFR-06	Performance	A catalogue read at 500 listings shall complete within 100 ms of application time.	Measurement of 200 repetitions per operation (section 6.9)
NFR-07	Performance	The cost of a catalogue read shall grow sub-linearly with the number of listings over the range 100 to 2,000 rows.	Scaling measurement (section 6.10)
NFR-08	Reliability	No request shall produce an unhandled server error for any	Suite of 80 automated cases including malformed and

ID	Attribute	Requirement	Verification
		input reachable through the interface or the API.	missing input (section 6.3)
NFR-09	Security	No password shall be stored or transmitted in recoverable form; credentials shall be stored as a salted hash.	Inspection of the stored column; assertion in test TC-A10
NFR-10	Security	All user-supplied text shall be escaped before it reaches the document so that a listing or comment cannot execute script in another user's browser.	Code inspection of every render path; adversarial input in a listing title
NFR-11	Security	Authorisation shall be enforced on the server for every state-changing operation, and shall not rely on the client hiding a control.	Authorisation test matrix (section 6.7)
NFR-12	Maintainability and portability	The system shall install and start with a single command on a clean Windows or macOS machine with Python installed, and shall need no server, container or build step.	Clean-environment installation of the submitted code (section 6.13)

3.4 Requirement Prioritisation

Requirements were prioritised using MoSCoW so that, if the schedule slipped, the team would cut the right things. The rule applied was that a **must** requirement is one whose absence breaks the transaction loop; everything else could be deferred without the software ceasing to be a proof of concept.

Table 3.3 *Prioritisation of requirements*

Priority	Requirements	Rationale
Must (18)	FR-01 to FR-07, FR-09, FR-11 to FR-16, FR-19, FR-22 to FR-25, FR-28	Without these a student cannot list an item, find an item or complete a transaction. Removing any one breaks the loop.
Should (8)	FR-08, FR-10, FR-17, FR-18, FR-20, FR-21, FR-26, FR-27	The system is usable without them, but the experience is significantly poorer: sellers cannot correct a listing, and the public availability answer disappears.
Could (2)	FR-13 combined with FR-15 as a single filtered-sort control	A refinement of an existing capability rather than new capability.
Won't (this project)	Payment, delivery, reputation, notification, recommendation, multi-campus operation	Each was rejected in scope discussions. Payment and delivery add legal and operational obligations; reputation and recommendation need a volume of data the proof of concept will never have.

3.5 System Context

Figure 3.1 places the system in its environment. The boundary encloses the Flask application and the client code it serves; everything outside it is a human actor, the browser, the filesystem holding the database and the photographs, or the assessor. The context is deliberately narrow - the proof of concept has no external service dependency, which is what makes it installable with one command.

System context - who and what the PoC interacts with




Figure 3.1 System context: the actors, the browser and the local filesystem the proof of concept interacts with

3.6 Use Case Model

Figure 3.2 shows the eighteen use cases and the two actors that initiate them. The distinction between the actors carries real weight in the design: the unshaded use cases are reachable without a session, and the shaded ones are guarded both in the client - which redirects an anonymous visitor to the authentication view - and, independently, on the server. The visitor actor generalises to the registered student, so an authenticated student inherits every use case available to a visitor.




Figure 3.2 Use case diagram: eighteen use cases across two actors, with the visitor actor generalising to the registered student

Table 3.4 summarises the eighteen use cases and the requirement each one satisfies. The eight use cases that carry the most engineering risk are then described in full in section 3.7; the remaining ten follow the same pattern and are recorded in the requirements traceability matrix in Appendix B.

Table 3.4 Use case summary

ID	Use case	Actor	Requirements
UC-01	Register an account	Visitor	FR-01
UC-02	Log in and log out	Registered student	FR-02, FR-03
UC-03	Browse the catalogue	Visitor	FR-11
UC-04	Search, filter and sort listings	Visitor	FR-12 to FR-15
UC-05	View a listing in detail	Visitor	FR-16
UC-06	Publish a listing	Registered student	FR-05, FR-07
UC-07	Upload listing photographs	Registered student	FR-06
UC-08	Edit an own listing	Registered student	FR-08
UC-09	Take a listing offline or relist it	Registered student	FR-09, FR-10
UC-10	Read and post a comment	Visitor, registered student	FR-17, FR-18
UC-11	Save or unsave a favourite	Registered student	FR-19 to FR-21
UC-12	Open the personal centre	Registered student	FR-04, FR-28
UC-13	Reserve an item	Registered student (buyer)	FR-22
UC-14	Confirm the sale	Registered student (seller)	FR-23
UC-15	Cancel a reservation	Registered student (seller)	FR-24
UC-16	Send a private message	Registered student	FR-25
UC-17	List conversations	Registered student	FR-26
UC-18	Open and read a conversation	Registered student	FR-27

3.7 Detailed Use Case Descriptions

The eight use cases below are given in full because each exercises logic that cannot be inferred from the requirement statement alone: validation with two enforcement points, a file-handling path, a state machine with two independent guards, and a query whose correctness depends on how a conversation is identified.

UC-06 Publish a listing

Actor	Registered student (seller)
Goal	The student has something to sell and wants it visible in the catalogue.
Preconditions	An authenticated session exists and the student is on the “Sell an item” view.
Main success scenario	1. The student opens the publishing view. 2. The system presents the form fields with the required ones marked, and a live preview of the listing card. 3. The student enters a title, a description and a price, and selects a category and a condition. 4. The student selects between one and three photographs. 5. The system validates, in the browser, that a title is present, that the price parses as a number not less than zero, and that at least one photograph has been chosen. Invalid input produces an inline message and the flow stops. 6. The system uploads the photographs and receives their stored names. 7. The system creates the listing with the status on sale and returns its identifier. 8. The system shows a confirmation, clears the form and returns the student to the catalogue with the new listing present.
Alternative and exception flows	5a. A photograph exceeds 5 MB or is not one of the accepted image types: the system names the offending file and does not add it. 5b. More than three photographs are selected: the system refuses the excess and explains the limit. 7a. The server-side validation rejects the request: the system reports the reason returned by the server and keeps the entered values so nothing is lost.
Requirements satisfied	FR-05, FR-06, FR-07

The validation in steps 5 and 7a is deliberately duplicated. The browser exists to give immediate feedback; the server exists to be authoritative, because a request can be constructed without the browser.

UC-07 Upload listing photographs

Actor	Registered student (seller)
Goal	The seller needs photographs of the item to be stored so a listing can refer to them.
Preconditions	An authenticated session exists and the seller has selected one or more files.
Main success scenario	1. The system receives the selected files. 2. The system checks each file's extension against the allow-list of image types and rejects the whole request if any file is not permitted. 3. The system derives an extension from the sanitised file name and generates a random

	unique name for each file. 4. The system writes each file into the upload directory. 5. The system returns the list of generated names. 6. The seller's client passes those names to the listing creation request, which records one row per photograph.
Alternative and exception flows	1a. No file is present in the request: the system returns a validation error. 2a. A file's extension is not permitted: the system returns an error naming the allow-list. 3a. The sanitised name has no extension: the file is treated as unsupported.
Requirements satisfied	FR-06

The generated name is what protects the upload directory: a name supplied by the client is never used as a path.

UC-13 Reserve an item

Actor	Registered student, acting as a buyer
Goal	The buyer wants an item held for them so it is not sold to someone else.
Preconditions	An authenticated session exists, the item exists, and the buyer is not the seller.
Main success scenario	1. The buyer asks to reserve the item. 2. The system confirms the intent, explaining that the seller will be notified. 3. The system looks up the reservation rule and reads the item's current record. 4. The system checks that the caller differs from the seller; if the caller is the seller, the request is refused. 5. The system checks that the stored status is on sale; if it is anything else the request is refused and the row is not modified. 6. The system writes the new status and the modification time. 7. The system confirms, and the detail view re-renders with the reserved state and the buyer's actions replaced by seller-only actions.
Alternative and exception flows	3a. The item does not exist: the system reports that it was not found. 4a. The buyer attempts to reserve their own listing: refused. 5a. The item is already reserved or sold: refused, with the current status named so the buyer understands why. 2a. The buyer dismisses the confirmation: no request is sent and nothing changes.
Requirements satisfied	FR-22

The two checks are independent and both are necessary. The role check establishes who may perform the action; the state check establishes when. A design with only the role check allows a reservation to overwrite a completed sale.

UC-14 Confirm the sale

Actor	Registered student, acting as a seller
Goal	The seller has handed the item over and wants the listing closed as sold.
Preconditions	An authenticated session exists and the caller is the seller of the listing.
Main success scenario	1. The seller asks to mark the item sold. 2. The system confirms the intent. 3. The system looks up the confirmation rule, reads the item's record and checks that the

	caller is the seller and that the status is reserved. 4. The system writes the sold status and the modification time. 5. The system confirms, and the listing leaves the on-sale catalogue and appears under the seller's sold items.
Alternative and exception flows	3a. The status is not reserved - for example the seller is trying to sell an item nobody reserved: refused. 3b. The caller is not the seller: refused with a permission error. 4a. Any further action on a sold item: refused, because sold has no outgoing transition.
Requirements satisfied	FR-23

The requirement that a sale may only be confirmed from the reserved state is what makes the status field trustworthy. Without it, a seller could mark an item sold that a buyer had put effort into reserving, and a buyer could not rely on the reservation meaning anything.

UC-15 Cancel a reservation

Actor	Registered student, acting as a seller
Goal	The buyer did not complete the handover and the seller wants the item available again.
Preconditions	An authenticated session exists, the caller is the seller, and the status is reserved.
Main success scenario	1. The seller asks to cancel the reservation. 2. The system confirms the intent, explaining that the item returns to sale. 3. The system checks the role and the current status. 4. The system writes the on-sale status and the modification time. 5. The item reappears in the catalogue and becomes reservable by any buyer.
Alternative and exception flows	3a. The status is not reserved: refused. 3b. The caller is not the seller: refused.
Requirements satisfied	FR-24

This is the recovery path that makes the reservation safe to use: a reservation that could never be released would be worse than none, because one unreliable buyer would remove the item from circulation permanently.

UC-11 Save or unsave a favourite

Actor	Registered student
Goal	The student wants to keep a record of an item they are considering, without committing to it.
Preconditions	An authenticated session exists and the item exists.
Main success scenario	1. The student asks to save the item. 2. The system records the association between the student and the item. 3. The system confirms, and the control changes to indicate the saved state. 4. The student later asks to unsave it; the system removes the association. 5. The student's favourites list, most recently saved first, is available from the personal centre.
Alternative and exception flows	1a. The item is already saved: the system refuses the duplicate rather than creating a second row, and reports that it is already saved.

	1b. No session exists: the request is refused.
Requirements satisfied	FR-19, FR-20, FR-21

Uniqueness is enforced by the database rather than by a read-then-write check, which would have a race between the read and the write.

UC-18 Open and read a conversation

Actor	Registered student
Goal	The student wants to continue a negotiation about an item.
Preconditions	An authenticated session exists and a conversation with the other student exists or is being started.
Main success scenario	1. The student opens the conversation with another student. 2. The system retrieves every message between the two students, in either direction, ordered by time. 3. The system marks the messages the student has received from the other party as read. 4. The system presents the messages as bubbles, distinguishing the student's own messages from the other party's, and scrolls to the most recent. 5. The student types a reply and sends it. 6. The system records the message and returns the student to step 2 so the conversation is re-rendered from the stored state.
Alternative and exception flows	1a. No conversation exists: the system presents an empty thread with an invitation to start one. 5a. The student sends an empty message: the system ignores the submission. 5b. The message cannot be stored: the system reports the failure and preserves what the student typed.
Requirements satisfied	FR-27

Marking messages read as part of reading them removes a class of bug in which the client forgets to report that something was read.

3.8 Business Process Model

Figure 3.3 models the process that the software supports, end to end. The three lanes separate the two human roles from the system, which is what makes the hand-off points visible: the transaction only leaves the buyer's lane twice, once when the reservation must reach the seller and once when payment and handover happen offline. The final lane carries the guards, because the guards are the system's contribution to the process - without them the process is exactly the informal one that already fails.

Business process - one complete second-hand transaction

One lane per role; time flows left to right. Dashed edges cross a lane boundary.




Figure 3.3 Business process model for a complete second-hand transaction, with the server-side guards shown in the system lane

3.9 Activity Model for Publishing

Figure 3.4 models publishing as an activity diagram, because it is the one process in the system with a genuine sequence of dependent decisions. Three decisions appear, and each is checked twice: once in the browser for immediate feedback and once on the server because the browser cannot be trusted. The diagram also makes visible why listing creation and photograph upload are separate requests: the upload completes before the listing is created, so a failure in the second step cannot leave a listing that refers to photographs which do not exist.

Activity diagram - publish a listing
Rounded bars are decisions, the guard is written on the outgoing edge.




Figure 3.4 Activity model for publishing a listing, showing the three decisions and the duplicated client-side and server-side enforcement

3.10 Requirements Traceability

Traceability is maintained in both directions. Forwards, every requirement names the use case that exercises it and the design element that implements it; backwards, every test case names the requirement it verifies, so a failing test identifies which requirement is now at risk as well as what broke.

4 System Design and Architecture

4.1 Architectural Style and the Alternatives Considered

The system is a **three-tier layered web application**: a framework-free client that runs in the browser, a Flask application that owns all logic and all authorisation, and a thin persistence layer holding a single SQLite file and an upload directory. Each layer talks only to the layer beneath it, and the client never reaches the database directly.

The style was chosen against three alternatives, and the deciding factor in each case was constraint C1 from section 2.5 - the software must run on the assessor's machine with one command and no configuration.

Table 4.1 *Architectural alternatives and the reason each was rejected*

Alternative	Attraction	Reason for rejection
Three-tier layered (chosen)	Separates presentation from logic and logic from storage; each layer can be changed independently; the whole system runs as one process with no external dependency	Accepted. It satisfies C1 while still giving a clean separation of concerns and an obvious place for authorisation.
Server-rendered pages with a template engine	Simpler than a client-side application: no JSON contract, no duplicated validation	Rejected because the interface needs to update without a full page load - the comment thread, the favourite control and the chat panel all change in place. Server rendering would have meant a page reload for each one.
Single-page client with a separate API service	The clearest separation of concerns, and the most common production shape	Rejected because it requires either a second process or a cross-origin configuration. Both break C1, and cross-origin requests would additionally require the session cookie to be sent explicitly.
Microservices	Independent deployment of accounts, listings and messaging	Rejected as disproportionate. The whole application is roughly 1,700 lines; splitting it would add service discovery, inter-service contracts and distributed failure modes to solve a scaling problem the system does not have.

4.2 Layered Architecture

Figure 4.1 shows the three layers and the responsibility of each. The boundary that matters most is the one between the presentation layer and the application layer: the client holds no authority. Every rule that protects data - who may change what, what the status is allowed to become next, which file types may be stored - is enforced below the boundary. The client knows those rules only in order to give better feedback, and the design accepts the cost of writing some validation twice in exchange for a system that cannot be subverted by constructing a request by hand.

Layered architecture and the responsibility of each layer




Figure 4.1 Layered architecture, showing the endpoint groups of the application layer and the components of the data layer

Two properties of this arrangement carry most of the design's weight. The presentation layer never reaches the database directly - every read and every write passes through a JSON endpoint - which means the application layer is the only place authorisation has to be correct. And the data layer is deliberately thin: it knows about tables and connections and nothing about rules. That thinness is what would make a future move to a hosted database a local change rather than a rewrite, and it is the reason the layer was kept thin even though the proof of concept does not need to be portable.

Table 4.2 Responsibility of each layer, and what each layer is forbidden to do

Layer	Owns	Must not
Presentation	View state, rendering, event handling, immediate input validation, escaping of all user-supplied text before it is inserted into the document	Hold authority over any data; construct SQL; be the only place a rule is enforced
Application	HTTP routing, session handling, authorisation, validation, the status state machine, construction of the JSON response envelope, file storage	Contain presentation markup; assume that a request originated from the provided client
Data	Connection creation, the relational schema, uniqueness constraints, file storage on disk	Contain business rules that the application layer also needs to know about

4.3 Module Decomposition

Figure 4.2 shows the modules and the direction of each dependency. The decomposition follows Parnas's criterion of hiding the decisions most likely to change (Parnas, 1972): the

shape of the data model is hidden behind the data layer, the shape of the HTTP contract is hidden behind the client's request wrapper, and the specifics of password hashing and file-name sanitisation are hidden behind a third-party library.

Module decomposition and the direction of each dependency




Figure 4.2 Module decomposition and the direction of each dependency

Table 4.3 Module inventory

Module	Lines	Responsibility	Depends on
app.py	552	The Flask application object and all 23 routes: accounts, catalogue, listing management, upload, transactions, favourites, comments, messaging and the personal centre	models.py, Werkzeug
models.py	106	Connection factory and schema creation for the six tables	sqlite3
templates/index.html	441	The single document: five view containers, two modal surfaces, an inline SVG icon sprite and the mount points for client rendering	style.css, api.js, main.js
static/js/api.js	63	One method per endpoint, so no URL is written twice and the request shape is defined in one place	the HTTP endpoints
static/js/main.js	1169	All view state, rendering, event binding, the hash router, escaping and the feedback layer	api.js
static/css/style.css	1025	Design tokens, component styles, responsive breakpoints	nothing

4.4 Domain Model

Figure 4.3 presents the domain model. It is described as persistence-oriented because the system was built without an object-relational mapper: the classes in the diagram are realised as tables and as dictionaries returned from queries, not as Python classes. Modelling them this

way was still worthwhile, because the association cardinalities are exactly what the schema must enforce.

Domain class diagram (persistence-oriented)




Figure 4.3 Domain class diagram, showing the six entities, their attributes and their associations

4.5 Database Design

The schema has six tables, shown in Figure 4.4. It is in third normal form: every non-key attribute depends on the whole key and on nothing else. The design was checked against the normal forms deliberately, and one candidate denormalisation - storing a reference to the cover photograph on the item row - was rejected in favour of a correlated sub-query, because the sub-query is trivially correct and a duplicated column would have needed code to keep it in step with the images table.

Entity-relationship diagram of the SQLite schema




Figure 4.4 Entity-relationship diagram of the six-table schema, with foreign keys and cardinality

Table 4.4 Table specifications

Table	Purpose	Key	Constraints and notes
users	One row per registered student	id	username is unique and not null; password_hash is not null; nickname, email, student_id and avatar are optional
items	One row per listing; the aggregate root	id	seller_id references users; title not null; price not null and real; status defaults to on sale; updated_at is written on every change
images	One row per photograph attached to a listing	id	item_id references items; file_path not null; the row stores the generated file name, never a client-supplied one
favorites	The many-to-many association between a student and a saved listing	id	user_id and item_id both reference their tables; UNIQUE(user_id, item_id) prevents a duplicate save
comments	A public question or answer attached to a listing	id	item_id and user_id both reference their tables; content not null
messages	A private message between two students, optionally about an item	id	sender_id and receiver_id both reference users; item_id references items and is nullable; is_read defaults to zero

Two properties of SQLite shaped the design and both are recorded as limitations rather than as oversights. The engine enforces declared foreign keys **only** when the connection is opened with `PRAGMA foreign_keys = ON`, and the proof of concept does not set that pragma; referential integrity is therefore currently a convention held up by the application layer. And `CURRENT_TIMESTAMP` resolves to whole seconds, which makes any ordering that relies on it alone unstable when two rows are written in the same second - a defect that the test suite found and that is reported in section 6.13.

4.6 Transaction State Machine

The status field of a listing is modelled as an explicit state machine rather than as a free-text field that the interface happens to set. Figure 4.5 shows the four states and the five transitions, and Table 4.5 sets out the rules. The model is the single most important design decision in the project, because it is what makes the answer to "is this still available?" authoritative rather than advisory.

State machine governing the life cycle of a listing




Figure 4.5 State machine governing the life cycle of a listing, with the actor and the guard for each transition

Table 4.5 Status transition rules. The role and the source state are both required; a request that fails either check changes nothing.

Action	From	To	Permitted actor	Failure response
reserve	on sale	reserved	Any student except the seller	403 if the caller is the seller; 400 if the status is not on sale
sell	reserved	sold	The seller	403 if the caller is not the seller; 400 if the status is not reserved
cancel_reserve	reserved	on sale	The seller	403 if the caller is not the seller; 400 if the status is not reserved
off_shelf	on sale	off shelf	The seller	403 if the caller is not the seller; 400 if the status is not on sale
relist	off shelf	on sale	The seller	403 if the caller is not the seller; 400 if the status is not off shelf
(any other)	-	-	-	400 invalid action, before the item is even read

Sold has no outgoing transition, which is a deliberate choice rather than an omission: undeleting a completed sale would require tracking which buyer completed it and by what

arrangement, and the proof of concept does not model the buyer side of a completed transaction at all. The consequence is that a mistaken confirmation cannot be undone through the interface, which is why the confirmation dialog for that action states explicitly that the listing will be marked as sold.

4.7 Interaction Design

Three interactions carry most of the system's behaviour and are modelled in full. Between them they exercise every layer, both directions of the status machine, and the one query in the system whose correctness is not obvious from reading it.

4.7.1 Publishing a listing

Figure 4.6 follows a publish action from the form submission to the committed rows. The sequence makes two design decisions explicit. The upload and the creation are separate requests, so the upload completes and returns file names before any listing row is written; a failure in the second step therefore cannot leave a listing that refers to a photograph that was never stored. And the listing row is written before the photograph rows, so that the generated key is available to be used as their foreign key - the reverse order would require a second read.

Sequence diagram - publish a listing, from form submit to committed row
Solid arrows are calls; dashed arrows are returns. Self-calls are work done inside one component.




Figure 4.6 Sequence diagram for publishing a listing, from form submission to committed rows

4.7.2 Reserving an item and confirming the sale

Figure 4.7 shows the two ends of the transaction. Both requests reach the same endpoint and are validated against the same rule table, which is the property that makes the state machine trustworthy: the two transitions cannot drift apart, because there is only one place where either is defined. The diagram also shows the failure path explicitly, because the behaviour of a rejected request - that the row is left untouched - is part of the specification rather than an implementation detail.

Sequence diagram - reserving an item and confirming the sale
Both transitions are validated against the same rule table before either row is written.




Figure 4.7 Sequence diagram for reserving an item and confirming the sale, including the two guards and the rejection path

4.7.3 Conversations and messaging

Figure 4.8 covers the three messaging interactions. The conversation list is produced by a single aggregate query rather than one query per conversation, which matters because a student's conversation list grows without bound while the number of students they are talking to does not. Opening a thread both reads and writes: the messages are retrieved and the incoming ones are marked as read in the same request, so unread state cannot get out of step with what the student has actually seen.

Sequence diagram - conversation list, opening a thread and sending a reply
The list is produced by one aggregate query rather than one query per conversation.




Figure 4.8 *Sequence diagram for the conversation list, opening a thread and sending a reply*

4.8 Interface Design

The client interacts with the server through a small JSON interface. Figure 4.9 groups the endpoints by the resource they act on, and Table 4.6 lists them with the methods they accept and whether a session is required. Every endpoint answers with the same envelope, described in Table 4.7, so that the client can handle success and failure uniformly without inspecting the shape of the body.

JSON endpoint surface, grouped by resource
 23 routes in total: 2 page routes, 21 JSON endpoints. Shaded badges mark methods that need a session.

Accounts			
POST	/api/register		public
POST	/api/login		public
POST	/api/logout		public
GET	/api/user/me		session
Catalogue			
GET	/api/items		public
GET	/api/items/<id>		public
Listing management			
POST	/api/items		session
PUT	/api/items/<id>		session
DELETE	/api/items/<id>		session
Media			
POST	/api/upload		session
GET	/uploads/<filename>		public
Transactions			
POST	/api/items/<id>/status		session
Social			
GET	/api/favorites		session
POST	/api/favorites		session
DELETE	/api/favorites/<id>		session
GET	/api/items/<id>/comments		public
POST	/api/items/<id>/comments		session
GET	/api/messages		session
POST	/api/messages		session
GET	/api/messages/<uid>		session
Personal centre			
GET	/api/user/items		session
GET	/api/user/sold		session

Figure 4.9 The JSON endpoint surface, grouped by resource, with the methods and the session requirement of each endpoint

Table 4.6 Endpoint catalogue. Session means that the endpoint refuses an unauthenticated request with 401 without evaluating the body.

Method	Path	Purpose	Session
POST	/api/register	Create an account	No
POST	/api/login	Authenticate and establish a session	No
POST	/api/logout	Terminate the session	No
GET	/api/user/me	Return the caller's own profile	Yes
GET	/api/items	Catalogue with keyword, category, status, sort and page	No
GET	/api/items/<id>	Listing detail with images and comments	No
POST	/api/items	Create a listing	Yes
PUT	/api/items/<id>	Edit a listing the caller owns	Yes
DELETE	/api/items/<id>	Take a listing offline (soft delete)	Yes
POST	/api/upload	Store one or more photographs and return their names	Yes
GET	/uploads/<name>	Serve a stored photograph	No
POST	/api/items/<id>/status	Apply a status transition	Yes
GET	/api/favorites	List the caller's saved listings	Yes
POST	/api/favorites	Save a listing	Yes
DELETE	/api/favorites/<id>	Unsave a listing	Yes
GET	/api/items/<id>/comments	Read the public comments on a listing	No
POST	/api/items/<id>/comments	Post a comment	Yes
GET	/api/messages	Conversation list, grouped by partner	Yes
POST	/api/messages	Send a private message	Yes
GET	/api/messages/<uid>	Read a thread and mark it read	Yes

Method	Path	Purpose	Session
GET	/api/user/items	The caller's own listings	Yes
GET	/api/user/sold	The caller's sold listings	Yes

Table 4.7 *The response envelope shared by every endpoint*

Situation	HTTP status	Body	Client behaviour
Success with a body	200	{ code: 200, data: ... }	Render the data
Success without a body	200	{ code: 200, msg: '...' }	Show the message as a confirmation
Validation failure	400	{ code: 400, msg: '<reason>' }	Show the reason as an error next to the offending control
Not authenticated	401	{ code: 401, msg: '...' }	Redirect to the authentication view and remember the intent
Not permitted	403	{ code: 403, msg: '<reason>' }	Show the reason; do not retry
Not found	404	{ code: 404, msg: '...' }	Show an empty state with a way back
Payload too large	413	HTML error page from the framework	Report that the file exceeds the 5 MB limit

4.9 Client-Side Design

The client is a single document with five view containers that are shown and hidden by a small router, shown in Figure 4.10. Three design decisions distinguish it. The **hash route** `#item-<id>` makes an individual listing addressable, so a listing can be pasted into a message and will open directly on the detail view; the route is parsed both on load and on every change of the hash. The **guards are duplicated**: the client redirects an anonymous visitor away from the publishing and personal views, and the server independently refuses the same requests. And **hidden containers are not rebuilt**: switching views re-renders only the content that has changed, so a half-typed listing form survives a trip to the catalogue.

Client-side view router

One HTML document, five views, and a hash route that makes a single listing shareable.




Figure 4.10 The client-side view router: five views, the persistent header, and the guards on the views that require a session

4.10 Design Decisions and Trade-offs

Table 4.8 records the decisions that shaped the system, the alternative in each case, and the cost that was knowingly accepted. Recording the rejected alternative matters: a decision log that lists only what was chosen cannot be reviewed, because there is no way to tell whether the alternative was considered and rejected or simply never noticed.

Table 4.8 Design decision log

#	Decision	Alternative rejected	Trade-off accepted
D1	A single Flask process serving both the client and the API	A separate API service or a second process	The client and the server cannot be deployed or scaled independently. Accepted because C1 forbids a second process.
D2	SQLite as the store	PostgreSQL or a hosted relational service	No concurrent writers, no network access for other clients, and foreign keys unenforced by default. Accepted because the schema then travels with the submission and nothing has to be installed.
D3	No object-relational mapper	SQLAlchemy or an equivalent	All 29 statements are written by hand and the schema shape is visible in the route handlers. Accepted because the query set is small and the aggregate in section 5.7 is easier to express directly than through a mapper.
D4	Validation duplicated in the browser and on the server	Server-only validation	Some rules are written twice, and the two copies can drift. Accepted because client-only validation is not validation at all, and server-only validation gives the user a round trip before any feedback.
D5	A rule table driving the status transitions	A chain of conditional statements per action	The legal transitions are described in a data structure rather than in control flow, which some readers find less direct. Accepted because the rule set is readable in one place and cannot disagree with the enforcement.
D6	Soft deletion for listings	Removing the row and its photographs	Offline listings remain in the database and their photographs are never reclaimed. Accepted because the seller's history and the buyer's conversation remain coherent, and because a mistaken deletion is recoverable.

#	Decision	Alternative rejected	Trade-off accepted
D7	Generated file names for uploads	Keeping the name supplied by the client	Uploaded files are not identifiable by name on disk. Accepted because a client-supplied name is a path-traversal risk and can collide with an existing file.
D8	Two independent guards on every transition	The role check alone	One extra read per request, and one more failure mode to test. Accepted because the role check answers who and the state check answers when, and a reservation that overwrote a completed sale would be a correctness failure.
D9	A single aggregate query for the conversation list	One query per conversation, or caching the list	The query is harder to read than a simple select. Accepted because it keeps the cost constant as the number of conversations grows, and the alternative is a query per row of the list being rendered.
D10	No payment capability	An integrated payment provider	The platform cannot take a commission, and handover must be arranged between the two students. Accepted deliberately: payment would bring financial regulation, dispute resolution and fraud handling into a course project.

5 Implementation

5.1 Technology Stack

The proof of concept is written in Python with Flask, uses SQLite for storage, and serves a client built from plain HTML, CSS and JavaScript with no framework and no build step.

Figure 5.1 gives the reasoning behind each choice; every component was selected against one test - can the assessor obtain and run it without a paid account, a container runtime or a compilation step?

Technology choices and the reasoning behind them

Every component was chosen to keep the proof of concept installable on a laptop with one command.

Python 3.8+ language	already available on the lab and staff machines; no runtime purchase
Flask 2.3.3 web framework	routing, sessions and JSON helpers in one small dependency
Werkzeug 2.3.7 security utilities	password hashing and secure_filename without extra libraries
SQLite 3 database	file-based, no server to install; the schema travels with the submission
Vanilla JS / CSS client	no build step, no node_modules; the marker opens index.html from Flask
Playwright + Edge verification	used to capture the interface states reproduced in this report

Figure 5.1 *Technology choices and the reasoning behind each one*

Table 5.1 *Component versions and roles*

Component	Version	Role in the system
Python	3.8 or later (developed on 3.13)	Runtime for the application and the test suite
Flask	2.3.3	Routing, request parsing, signed sessions and JSON responses
Werkzeug	2.3.7	Password hashing and file-name sanitisation
SQLite	3 (bundled with Python)	Relational storage, one file on disk
pytest	9.x	Test runner used for the suite reported in Chapter 6
Playwright (Node)	via playwright-core	Drives Microsoft Edge for the interface verification and the screenshots in this report

5.2 Software Structure and Size

Table 5.2 gives the size of each component of the delivered software. The distribution is worth noting: the client is longer than the server, which is the expected shape for a system whose complexity lies in interaction rather than in computation. The server is 658 lines for 21 JSON endpoints, an average of 31 lines per endpoint including validation and authorisation.

Table 5.2 *Size of each component of the delivered software*

Component	Lines	Share
static/js/main.js	1,169	34%
static/css/style.css	1,025	30%
app.py	552	16%
templates/index.html	441	13%
models.py	106	3%
static/js/api.js	63	2%
tests/ (test suite and harnesses)	~1,300	not part of the runtime
Total runtime code	3,356	100%

5.3 Accounts, Sessions and Credential Storage

Registration accepts a username, a password and three optional profile fields. The password is never stored: it is passed through a salted key-derivation function and only the resulting hash is written to the database. The implementation deliberately uses the library's default parameters rather than choosing its own, on the principle that a hand-rolled KDF is a liability (OWASP, 2024). A test asserts that the stored value does not contain the submitted password and that it carries the expected algorithm prefix, so a future change that accidentally stored the plaintext would fail the suite rather than pass silently.

Sessions are held in Flask's signed session cookie. The application's only interaction with them is to write the authenticated user's identifier on login and to remove it on logout. A decorator applied to every protected route reads that identifier and refuses the request with 401 if it is absent. Because the decorator is applied at the route level rather than checked inside handlers, a new endpoint added without it stands out immediately in review - and the authorisation matrix in section 6.7 exists to make that failure mode impossible to miss.

5.4 The Catalogue Query

The catalogue endpoint accepts five optional parameters - keyword, category, status, sort and page - and builds one statement from them. The construction is done by appending fragments to a list of conditions and extending a parallel list of bound parameters, never by interpolating values into the statement text. The distinction matters: interpolating a keyword would make the search box an injection point, whereas binding it cannot. The same condition list is reused to build the counting query, so the total that drives the pagination control is computed from exactly the same filter as the page it describes.

The cost of the wildcard search is a full scan of the items table, because a leading wildcard cannot use an index. At the scale the proof of concept operates on this is irrelevant, and

section 6.10 measures it; at a larger scale the answer is a full-text index rather than a change to the query.

5.5 Photograph Upload and Storage

Upload is the one part of the system that touches the filesystem, and it is treated accordingly. The endpoint accepts a multipart request containing one or more files, checks each file's extension against an allow-list, derives the extension from a sanitised copy of the supplied name, and then writes the file under a name generated from a random universally unique identifier. The name supplied by the client is used for nothing except to read its extension. Table 5.3 lists the rules and where each is enforced.

Table 5.3 *Validation rules applied to an upload or a listing, and where each is enforced*

Rule	Limit	Enforced in the browser	Enforced on the server
Accepted file types	png, jpg, jpeg, gif, webp	By MIME type at selection	By extension against the allow-list
Maximum size per file	5 MB	At selection	By the framework's request size limit (413)
Maximum files per listing	3	At selection	Not enforced; the client is the only place the limit is applied
Title required	Non-empty	Before submission	Before the row is written
Price required and numeric	Not negative	Before submission	Before the row is written
At least one photograph	One	Before submission	Not enforced; a listing may legitimately have no photograph

The last two rows of Table 5.3 are the honest gaps. The three-photograph cap and the requirement for at least one photograph exist only in the client. A request constructed by hand can publish a listing with no photograph or with thirty. Neither omission corrupts data, and the interface never produces such a request, so the risk was accepted for the proof of concept and is listed as a defect to fix rather than presented as a complete implementation of the requirement.

5.6 The Status Machine

The status machine is implemented as a dictionary that maps each action to the required source state, the resulting state and the role permitted to perform it. The endpoint then performs the same five steps for every action, shown in Figure 5.2: reject an unknown action, load the item and reject a missing one, check the role, check the source state, and only then write. Because the steps are shared, adding a state would mean adding one entry to the table and one control to the interface, with no change to the logic that validates transitions.

How a status request is validated before the row is written

The same five-entry rule table drives every transition, so authorisation and state cannot drift apart.




Figure 5.2 How a status request is validated before the row is written, driven by a five-entry rule table

Two design choices in this component are worth defending. The action is validated against the table **before the item is read**, so an unknown action costs nothing and cannot reveal whether an item exists. And the write is the final step with no conditional logic after it, which is what guarantees the property the tests assert: a request that fails any check leaves the row exactly as it was.

5.7 The Conversation List Query

The conversation list is the one algorithm in the system that needed actual thought. The requirement is to show one row per person the student is talking to, containing the most recent message in that thread. The obvious implementation - iterate the student's messages, work out the other party for each, and keep the newest per person - costs one pass in the application and requires either a query per conversation or all messages loaded into memory. Figure 5.3 shows the single aggregate query that replaces it.

Rendering the conversation list without a query per conversation

One grouped aggregate returns the newest message of every thread the caller takes part in.




Figure 5.3 Rendering the conversation list with one aggregate query instead of a query per conversation

The query works in two parts. The inner part groups the student's messages by the other party, using a conditional expression to normalise the pair of sender and receiver into a single

identity regardless of direction, and selects the largest identifier in each group. The outer part joins back to the messages table on that identifier and to the users table to obtain the partner's display name. The cost is one pass plus one grouped sub-query, and it does not grow with the number of conversations.

5.8 Client-Side Implementation

The client is one document containing five view containers, two modal surfaces, an inline icon sprite and the mount points that the rendering code fills. All rendering is driven from state held in module-level variables, and every path that inserts user-supplied text into the document passes it through an escaping function first. That function is the system's principal defence against cross-site scripting, and it is applied to titles, descriptions, comment bodies, message bodies, nicknames and file names. Table 5.4 lists the paths and what each escapes.

Table 5.4 *Client-side rendering paths and the escaping applied to each*

Rendering path	User-supplied values interpolated	Escaping
Listing card	title, seller nickname, price, category	All values escaped; the price is coerced to a number before formatting
Listing detail	title, description, seller nickname and email, category, condition	All values escaped
Comment list	author nickname, body, relative time	All values escaped; times are formatted, never inserted raw
Chat and conversation list	message body, partner nickname	All values escaped
Publish preview	title, price, category	All values escaped; the preview is built by the same function as the catalogue card
Profile head	nickname, username, student identifier	All values escaped
Toast and confirm dialog	server-supplied message	Escaped for toasts; the confirm dialog assigns through <code>textContent</code> , which cannot create markup at all

5.9 Design System

The visual language is defined once, as custom properties on the document root, and every component style refers to those properties rather than to literal values. Table 5.5 lists the token groups. The practical benefit is that the whole theme can be retuned from one place, and the review benefit is that a component that hard-codes a colour stands out when the file is read.

Table 5.5 *Design token groups and what each controls*

Group	Tokens	Used for
Colour	Surface, surface-muted, border, ink, ink-muted, accent, accent-soft, success, danger, warning	Every background, border, text colour and status badge in the interface

Group	Tokens	Used for
Typography	Two sizes for body text, four for headings, one monospace family, a single line-height scale	All text; no component sets its own font size outside the scale
Spacing	A five-step scale from 4 px to 40 px	All margins, padding and grid gaps
Radius and elevation	Three radii and three shadow levels	Cards, buttons, dialogs and the sticky purchase panel
Breakpoints	1040 px, 900 px, 760 px	Collapsing the detail page to one column, the grid to two columns, and the header to a compact form

5.10 User Interface

The interface is presented as a sequence of states rather than as one screenshot per screen, because the states that carry the most design work are the ones that appear only when something is empty or when a request is refused. Every figure below was captured from the delivered software driving a real browser at a 1240-pixel viewport, and the capture run reported no console errors and no unexpected 4xx or 5xx responses. The figures are cropped to the region under discussion: a full-page capture reduced to the width of a printed text column leaves body type around five points, which is not readable, whereas a crop placed at a smaller physical width keeps it at eight or nine.

5.10.1 Browsing the catalogue

The catalogue is the entry point and has to work for a visitor with no account. Figure 5.4 shows the listing grid with its sort control. The cards carry the status badge, the cover photograph, the price, the seller and the category, which is the set of attributes the competitive review in section 2.2 identified as the ones a buyer screens on.




Figure 5.4 The catalogue grid: status badge, cover image, price, seller and category on each card

Figure 5.5 shows the empty state. It is included deliberately, because the empty state is where projects of this shape usually stop: the grid simply stays blank and the user cannot tell whether nothing matched or something failed. The implementation names the condition and offers a next step.




Figure 5.5 The empty search result: the condition is named and a next step offered, instead of leaving the grid blank

5.10.2 Reading a listing and acting on it

The detail view is the densest screen and the one where the status machine becomes visible to a user. Figures 5.6 and 5.7 show the same listing from the two sides of a transaction. For a signed-in buyer who is not the seller, the available actions are to reserve the item, to save it and to message the seller, and the safety panel states the platform's position on payment plainly. For the seller, the action set is entirely different: take the listing offline while it is on sale, or confirm the sale and cancel the reservation while it is reserved. The interface is role-aware, but the authority behind it is not - the server would refuse the same actions from the wrong account even if the controls were present, which is what the authorisation matrix in section 6.7 establishes.




Figure 5.6 *The detail view as a signed-in buyer sees it: reserve, save and message, with the seller and the structured attributes*




Figure 5.7 *The same view for the seller: the buying actions are replaced by the management actions*

Figure 5.8 shows the confirmation dialog that guards taking a listing offline. Every destructive or hard-to-reverse action is routed through this dialog, and it states what will happen rather than asking a bare "are you sure?"




Figure 5.8 *The shared confirmation dialog used for every destructive action in the interface*

5.10.3 Publishing

Figure 5.9 shows the publishing form. Three decisions are visible. The category and condition are chip selectors rather than dropdowns, kept in step with hidden select elements so the form remains submittable without script. The photograph control is a drop zone that also accepts a click, with thumbnails that can be removed individually. And a live preview of the listing card sits beside the form, built by the same function that renders the catalogue, so what the seller sees is what buyers will see.

Publish an item

More complete the listing, the faster it sells. Fields marked * are required.

Title *

Desktop monitor stand, aluminium

One line about what it is and its condition, up to 60 characters

Description

Used for one semester, no scratches. Height adjustable, fits a 27-inch screen.

Price *

¥ 45

Category

Textbook Electronics Daily use Sports gear Other

Condition

Brand new Like new Good Fair

Photos *

Click to choose photos, or drag them here

Figure 5.9 The publishing form: chip selectors for category and condition, a photograph drop zone, and validation messages attached to the field that failed

5.10.4 The personal centre and messaging

The personal centre gathers everything belonging to the signed-in student behind four tabs. Figure 5.10 shows the profile head with its three summary counts and the own-listings tab; the saved listings and sold tabs reuse the same card component as the catalogue, which is why the grid looks identical in all three places.



Brian Li

@demo_li · Student ID 2023010233




Figure 5.10 *The personal centre: profile summary with counts, and the student's own listings*

Figure 5.11 shows an open conversation. Messages from the student and from the other party are aligned differently and labelled with relative times, the panel scrolls to the most recent message, and Enter sends. The panel is labelled a direct message rather than a live chat, because the implementation fetches on opening and after sending rather than polling - a deliberate choice, since polling would spend the system's complexity budget on a feature the evaluation does not require. The relative time is computed in the browser from the stored timestamp, and the parsing normalises the value to UTC explicitly first, because the database stores timestamps without a zone and a naive parse produces an error equal to the machine's offset from UTC.




Figure 5.11 *An open conversation: own and other messages distinguished, relative times, and the composer anchored to the bottom*

5.10.5 Small screens

Figure 5.12 shows the interface at a 400-pixel viewport, a typical large-phone width. The grid collapses to two columns, the header reduces to the brand, the search field and a compact action, and the detail page collapses to a single column with the purchase panel below the gallery. This is the viewport used for requirement NFR-05.




Figure 5.12 *The interface at a 400-pixel viewport: the grid collapses to two columns and the header reduces to its essential controls*

5.11 Coding Standards and Conventions

The code follows a small set of conventions that were agreed before implementation and checked in review. They are recorded here because they are part of what the software engineering assessment examines, and because a convention that is not written down cannot be enforced.

Table 5.6 *Coding conventions applied across the codebase*

Convention	Rationale
Every SQL statement binds its parameters; no value is ever interpolated into statement text	Removes injection as a class of defect rather than defending against it case by case
Every response uses the same { <code>code</code> , <code>msg</code> , <code>data</code> } envelope	The client can handle success and failure in one place instead of inspecting bodies
Every protected route carries the same authorisation decorator	A missing guard is visible at the route definition rather than hidden inside a handler
Database connections are opened and closed inside the request that uses them, and closed on every path including the error paths	SQLite is file-based, so a leaked connection holds a lock rather than a socket
Client code escapes every interpolated value through one shared function	One place to get right, and a reviewable list of what is escaped
No CSS value is written outside the token set except where a token would be meaningless	Keeps the visual language consistent and makes a deviation visible
No SQL and no HTML is constructed by string concatenation in the data layer	Keeps the layering honest: the data layer knows about tables, not about markup

5.12 Deployment and Reproducibility

The application starts with a single command, `python app.py`, on any machine with Python 3.8 or later and the two declared dependencies installed. On first run it creates the database file and the upload directory, so a clean checkout becomes a running system without a migration or a seeding step. The repository ships with a demonstration dataset so every screen has content, and with a backup of the empty database so the data can be removed without editing the schema.

6 Testing and Evaluation

6.1 Test Strategy

Testing was organised to answer one question at each level: at the bottom, is this unit of logic correct in isolation; in the middle, do the parts work together across the interface between them; and at the top, does the assembled system behave correctly against the requirements, and acceptably under load. Table 6.1 sets out the levels, the technique used at each and what each level is able to detect.

Table 6.1 *Test levels, techniques and detection capability*

Level	Technique	What it detects	What it cannot detect
Unit and component	Automated cases against a Flask test client with a fresh database and upload directory per case	Incorrect logic in a single endpoint: validation, authorisation, state handling, query construction	Anything that depends on the assembled interface, and anything that depends on concurrency
Interface	Automated cases that assert the HTTP status and the response envelope for every endpoint, in both an authenticated and an anonymous context	Contract violations: a wrong status code, a missing guard, a body the client cannot parse	Whether the client interprets the contract correctly
Integration	Cases that exercise a user-level flow across several endpoints - publish, browse, reserve, confirm - against one database	Inconsistencies between endpoints, such as a listing that reports one status in the catalogue and another on its detail page	Rendering defects and usability problems
System and interface acceptance	Scripted browser runs over every screen in both authentication states, with console and network monitoring	Rendering failures, console errors, unexpected 4xx and 5xx responses, viewport breakage	Subjective usability; anything a real user would do that the script does not
Performance	Repeated timed measurement per operation, and a measurement of read cost as the dataset grows	Costs that exceed the stated budget, and growth that contradicts the stated scaling requirement	Behaviour under concurrent load, which a single-process SQLite deployment cannot expose meaningfully

6.2 Test Environment and Harness

The suite runs under pytest with a fixture that points the data layer at a temporary database file and the upload directory at a temporary directory, and creates the schema for each case. The consequence is that no case can be affected by another, and the repository's demonstration database is never touched. Because the fixture replaces the database path rather than the connection, the cases exercise the same code path that the running application uses; nothing is stubbed. The one exception is the framework's own request-size limit, which is exercised by sending a body larger than the configured maximum and asserting that the framework refuses it.

6.3 Unit and Component Test Results

The suite contains **80 cases**; all 80 pass. The cases are named and grouped by component so that a failure identifies which module regressed without needing to read the failure trace.

Figure 6.1 shows the distribution and Table 6.2 states what each group covers.

Automated test suite: 80 cases, all passing

Grouped by the component under test, so a failure names the module that regressed. Every case ran against a throwaway database and upload directory.




Figure 6.1 Automated test cases by component. All 80 cases pass.

Table 6.2 Automated test suite by group

Group	Cases	Range	Covers
TC-A Accounts	10	TC-A01 to TC-A10	Registration and its rejection cases, password hashing, login and its rejection cases, logout, profile retrieval, and the requirement that no credential material is returned
TC-C Catalogue	14	TC-C01 to TC-C14	The empty catalogue, public access, keyword matching against both searchable columns, case-insensitive matching, category and status filters, all three sort orders, pagination boundaries, detail assembly and a missing item
TC-L Listings	12	TC-L01 to TC-L12	Creation with each missing field, non-numeric price, the initial status, editing, ownership on edit and delete, the soft-delete behaviour, and the effect of an offline listing on the catalogue
TC-U Upload	8	TC-U01 to TC-U08	Absent session, empty request, accepted types, the generated name format, a disallowed extension, multiple files, an oversized body, and the order in which stored names are attached to a listing
TC-T Transitions	15	TC-T01 to TC-T15	Unknown and missing actions, a missing item, the session requirement, both guards on every action, the happy path of all five transitions, and the requirement that a failed guard leaves the row unchanged
TC-S Social	16	TC-S01 to TC-S16	Saving, unsaving and duplicate saves, comment ordering and public visibility, empty comments, sending and reading messages, the conversation grouping and its ordering behaviour, and the marking of messages as read
TC-P Personal centre	5	TC-P01 to TC-P05	Own listings only, the empty sold list, the sold list after a completed transaction, the cover-image key, and the session requirement

6.4 Interface Testing

Every endpoint is asserted on its HTTP status and on the shape of its envelope. Table 6.3, reproduced in Appendix C.1, records the status codes the interface is expected to return and confirms that each was observed. The table is a contract, not a description: a change that made any cell untrue would be a defect even if no user noticed it.

6.5 Negative and Boundary Testing

Negative testing is where most of the defects in a system of this shape are found, because the happy path is exercised by anyone who uses the software for five minutes. Table 6.4 lists the boundary and negative cases that were written deliberately, with the boundary being tested in each.

Table 6.4 *Negative and boundary cases written deliberately*

Case	Boundary under test	Expected and observed behaviour
TC-L03, TC-L04	Required field absent	400, with the reason naming the missing field
TC-L05	Non-numeric price	400; the value is rejected before any coercion is attempted
TC-C01	Zero rows in the table	The catalogue returns an empty list, a total of zero and zero pages rather than failing on a division in the pagination calculation
TC-C07	Keyword matching nothing	Empty result with a total of zero; the interface then shows the empty state of Figure 5.6
TC-C12	Pagination boundary at 21 listings, 12 per page	Page 1 holds twelve, page 2 holds nine, and the reported page count is two
TC-U03, TC-U05, TC-U06	File type allow-list at the boundary	Accepted types stored; a disallowed extension rejected for the whole request; multiple files in one request all stored
TC-U07	Body larger than 5 MB	413 from the framework, not a partial write
TC-A03, TC-A04	Duplicate and empty credentials	400 in both cases, before any row is written
TC-S02	Duplicate favourite	400 from the uniqueness constraint rather than a second row
TC-S06	Whitespace-only comment	400; the value is trimmed before the emptiness check, so a space is not accepted as content
TC-S10, TC-S11	Recipient that does not exist, and an empty body	404 and 400 respectively, both before any row is written
TC-T01, TC-T02	Unknown action, and an absent action field	400 before the item is read, so neither reveals whether an item exists
TC-T14	Actions on a terminal state	400 for relist, off-shelf and sell on a sold item; the state machine has no outgoing transition from sold

6.6 State Machine Testing

The status machine is the component where a defect would damage data rather than merely annoy a user, so it is tested exhaustively rather than representatively. The test design used a coverage criterion of **all transitions plus all illegal combinations**: for each of the five actions, the legal case plus a wrong-role case and a wrong-source-state case where those are

distinguishable. Table 6.5, reproduced in Appendix C.2, records the matrix with the outcome of each cell taken from the suite.

One case in this table is worth drawing out because it is the property the whole design exists to guarantee. TC-T15 takes a reserved item, switches to the buyer's session, and asks to confirm the sale. The request is refused with 403, and the test then reads the item back and asserts that its status is still reserved. A design that checked only the role, or that wrote the new status before validating, would fail this assertion while still returning the same status code to the caller. Testing the response alone would not have detected it.

6.7 Authorisation Testing

Authorisation is tested as a matrix rather than as a set of individual cases, because the failure mode that matters is not a wrong answer but a missing check. Table 6.6, reproduced in Appendix C.3, records for each protected operation whether it is refused for each of the three caller contexts. The pattern is what matters: every protected operation has a refusal in at least one column, and no operation is permitted to an anonymous caller.

The table exposes one asymmetry that is a deliberate design choice rather than a gap. Any signed-in student and any anonymous visitor may read every listing and every public comment; there is no ownership on read. That is what makes the catalogue a catalogue. Private messages are the exception, and the conversation read is scoped to the caller by the query itself rather than filtered afterwards, so a caller cannot retrieve another pair's messages by supplying their identifiers.

6.8 Requirements Coverage

Figure 6.2 summarises how many of the specified requirements are realised in the proof of concept and independently verified by at least one automated case. All twenty-eight functional and all twelve non-functional requirements are traced; the matrix that supports the figure is reproduced in Appendix B.

Requirements realised and independently verified

Each requirement below is traced to the design element that implements it and to at least one automated case that exercises it (Appendix B of the report).




Figure 6.2 *Requirements realised and independently verified, by category*

Coverage is complete at the level of the requirement, but it is not uniform in strength, and the difference matters when the results are read. Three requirements are verified more weakly than the rest. **FR-06**, which requires that up to three photographs be attached, is enforced only in the client: the server accepts an arbitrary number, so an automated case cannot demonstrate the cap. **NFR-04**, keyboard operability, is verified by a manual walkthrough rather than by an automated assertion, which means it is verified once rather than on every run. And **NFR-01**, that a first-time user can publish without instructions, rests on observation of three users rather than on a measurement - a small sample, reported as such.

6.9 Performance Testing

Performance was measured in-process with the framework's test client, which removes network and browser time and isolates the cost the application itself contributes. The dataset contains 500 listings, 22 accounts, 61 conversations and one 40-message thread; each read was repeated 200 times and each write 100 times. Table 6.7 and Figure 6.3 report the median and the 95th percentile per operation.

In-process response time by operation

Flask test client, 200 repeats per operation (100 for the write paths), Python 3.13.14. Network and browser time excluded.




Figure 6.3 In-process response time per operation at 500 listings, median and 95th percentile

Table 6.7 In-process response time per operation. Dataset: 500 listings, 22 accounts, 61 conversations, one 40-message thread. Python 3.13.14, 200 repetitions for reads and 100 for writes.

Operation	Median (ms)	95th pct (ms)	Mean (ms)	Repeats
Browse the catalogue (page 1 of 42)	2.55	2.73	2.59	200
Keyword search across title and description	2.58	2.67	2.60	200
Keyword search with no match	2.51	2.63	2.54	200
Filter by category and sort by price	2.51	2.72	2.54	200
Open a listing detail page	2.79	2.96	2.82	200
Open the busiest listing (5 comments)	2.84	3.05	2.88	200
Read the comment thread	1.89	2.08	1.92	200
Load the personal centre listing grid	3.49	3.92	3.54	200
Conversation list (60 threads)	1.95	2.07	1.97	200
Open a 40-message thread	2.44	2.68	2.69	200
Publish a listing	4.48	4.91	4.54	100
Add and remove a favourite	9.02	9.65	9.14	100
Toggle a listing offline and back	10.13	10.68	10.27	100

Three readings of these numbers matter. The catalogue reads sit between 1.9 and 3.5 ms of median application time, which is more than two orders of magnitude inside the 100 ms budget NFR-06 states, and the variation between the cheapest and the most expensive read is under 2 ms. The write paths cost roughly twice the reads, which is consistent with each write performing validation, a read for authorisation and a transaction commit. And the two largest

numbers in the table are not single operations at all: they are two-request cycles, so their per-request cost is half what the row shows.

6.10 Scaling Measurement

Absolute latency does not answer the question that matters for a catalogue: does the system degrade as the data grows? To measure it, the same three read paths were measured against datasets of 100, 250, 500, 1,000 and 2,000 listings, with the database file growing from 53 KB to 254 KB. Table 6.8 and Figure 6.4 report the medians.

How the catalogue reads scale with the size of the table

A 20-fold increase in rows costs about 0.3 ms. Both the keyword search and the filtered catalogue still scan the table, but at this scale the scan is not the dominant cost.




Figure 6.4 Median response time for three catalogue read paths as the number of listings grows from 100 to 2,000

Table 6.8 Median response time for three catalogue read paths as the dataset grows

Listings	Database size	Catalogue page 1 (ms)	Keyword search (ms)	Filter and sort (ms)
100	52 KB	2.45	2.47	2.43
250	68 KB	2.46	2.49	2.43
500	92 KB	2.56	2.58	2.53
1000	144 KB	2.55	2.64	2.52
2000	248 KB	2.75	2.86	2.63

The result is close to flat: over a twenty-fold increase in rows, the median read cost rose from about 2.45 ms to about 2.75 ms, an increase of roughly 12 per cent. Both the unfiltered catalogue and the keyword search scan the table, because a wildcard cannot use an index, yet the scan is not the dominant cost at this size - connection setup, row materialisation and JSON serialisation dominate. NFR-07 requires sub-linear growth and the measurement shows growth that is sub-linear by a wide margin, but the requirement is satisfied at a scale where the linear term has not yet become visible rather than because the algorithm is sub-linear. At

a scale where it did become visible, a wildcard search would need a full-text index; the measurement here cannot predict where that point falls, and does not claim to.

6.11 Interface and Usability Evaluation

The interface was evaluated in three ways, each of which finds different problems. A **scripted walkthrough** drove a real browser through every view in both an authenticated and an anonymous state, at a desktop and a mobile viewport, and captured the seventeen screens reproduced in section 5.10. The run monitored the browser console and every network response: the only reported console message in the whole run was the expected 401 that the client receives when it asks for the current user before anybody has logged in, and there were no unexpected 4xx or 5xx responses. That is a meaningful result, because a rendering path that throws inside an event handler leaves the screen looking correct while silently failing to respond to the next click.

Observation of three first-time users addressed requirement NFR-01. All three published a listing without written instructions. Two of the three attempted to publish without a photograph and were stopped by the client-side check with a message that named the problem; both then succeeded. That observation produced the only interface change to come out of the evaluation: the photograph control originally carried no indication that at least one image was required, and the field label was changed to mark it as required as a result.

6.12 Verification of the Non-Functional Requirements

Table 6.9, reproduced in Appendix C.4, states the measured outcome for each of the twelve non-functional requirements. Where the outcome is a partial pass, or where verification rests on a weaker method than the others, that is stated rather than averaged away.

6.13 Defects Found and Fixed

Table 6.10, reproduced in Appendix C.5, records the defects that were found and what happened to each. The entries are included because the way a defect was found says something about the testing that found it, and because two of them are still open - a report that lists only fixed defects is a report that stopped testing too early.

6.14 Threats to Validity and What Was Not Tested

The results above support the claims in this report, but they do not support more than those claims, and the boundary is worth stating precisely.

- **Concurrency was not tested.** The deployment is a single process with a file-based database. Concurrent writes are not exercised by any case, and the behaviour under concurrent load is not merely unmeasured but structurally different from the measured behaviour: SQLite serialises writers, so the timings in section 6.9 do not extrapolate.
- **The performance numbers exclude everything outside the application.** The measurement is in-process, so network latency, browser rendering, image transfer and the cost of the first request after the process starts are all absent. The numbers are a floor on user experience, not an estimate of it.
- **The usability evidence rests on three users.** Requirement NFR-01 was assessed by observation of three first-time users with no recordings and no measured task times. It is enough to have found one real interface problem; it is not enough to make a claim about learnability in general.
- **Accessibility was checked by walkthrough, not by audit.** Keyboard operability and the declaration of modal surfaces were verified manually. No screen-reader testing, no contrast measurement against a standard and no conformance audit against WCAG were performed, so the system should not be described as accessible, only as having taken specific accessibility steps.
- **Defensive testing has a surface, and the surface was drawn by the team.** The cases were written by the people who wrote the code. That is the weakest form of test design: a misunderstanding shared by the implementer and the tester produces a test that agrees with the misunderstanding. Independent test design, or an external review, would raise confidence and did not happen.
- **One platform was actually tested.** The software was developed and verified on Windows. The macOS claim in NFR-12 is reasoned from the absence of platform-specific code rather than observed.

Stating these limits is not a hedge on the results. Within the tested surface - the specified functional requirements, the twelve non-functional requirements, the transition guards and the authorisation rules - the evidence is strong and, in the case of the transition guards, exhaustive. Beyond that surface the report claims nothing.

7 Conclusions and Future Work

7.1 Outcomes Against the Objectives

Section 1.3 stated five objectives. Table 7.1 states, for each, what was achieved and what was not, judged against the evidence in Chapters 3 to 6 rather than against the intention.

Table 7.1 *Outcome against each project objective*

Objective	Achieved	Evidence	Shortfall
1. Specify the requirements	Yes	28 functional and 12 non-functional requirements, 18 use cases with 8 described in full, a prioritisation, a scope boundary and a traceability matrix (Chapters 3 and Appendix B)	Two requirements are implemented more weakly than specified: the photograph cap and the requirement for at least one photograph are enforced only in the client.
2. Design the architecture and data model	Yes	Architecture, module, domain, entity-relationship, state, activity and business process models, three sequence diagrams, and a ten-entry decision log with the rejected alternative recorded for each	The design assumes a single writer and does not address concurrency, which is acceptable at the stated scope but is a design gap rather than a delivery gap.
3. Implement a working proof of concept	Yes	3,356 lines of runtime code across 23 routes, implementing every feature listed in Table 1.3, running from a clean checkout with one command	Publishing a listing requires a photograph only in the client, and the interface does not poll for new messages.
4. Verify the proof of concept	Yes, with stated limits	80 automated cases all passing, a transition coverage matrix, a 16-operation authorisation matrix, a per-operation performance measurement and a five-size scaling measurement, plus an interface evaluation with no console errors and no unexpected error responses	Concurrency, real-user usability at scale and WCAG conformance were not tested. Two defects remain open and both are documented rather than hidden.
5. Deliver the documentation and presentation	In progress	This report, the requirements and design models in it, and the user manual, which is written against the delivered interface	The slide deck and the recorded demonstration were not complete at the time of writing.

7.2 Limitations

The limitations below are the ones a reviewer would find by reading the code and the tests. They are listed in descending order of how much they would matter to a real deployment, not of how hard they are to fix.

Table 7.2 *Limitations of the delivered proof of concept*

#	Limitation	Consequence if not addressed	Effort to fix
L1	The session signing key is a literal in the source file	Anyone with the source can forge a session cookie. Serious in any deployed system; limited here only because the source and the data are on the same machine.	Low: read it from the environment
L2	SQLite serialises writers and the service is single-process	The system cannot scale beyond one process, and the measured timings do not extrapolate to concurrent load.	High: a client-server database plus connection pooling
L3	Foreign keys are declared but not enforced, because the pragma is not set	Orphan rows are possible if a future code path deletes a referenced row. No current path does.	Very low: one statement on connection creation
L4	The conversation list orders by identifier rather than by timestamp	If rows ever arrive out of order the displayed "most recent" message can be wrong.	Very low: order by timestamp with the identifier as a tie-breaker

#	Limitation	Consequence if not addressed	Effort to fix
L5	The upload limits are partly client-side only	A hand-built request can exceed three photographs per listing, or create a listing with none.	Low: two validation checks on the server
L6	Uploaded photographs are stored on the local filesystem	The images and the database are not backed up together, and the application must run on the machine that holds them.	Moderate: object storage and a proxy route
L7	No mechanism to recover from a mistaken sale confirmation	A seller who confirms by accident cannot reverse it through the interface.	Moderate: an administrative path and a decision about what it is allowed to undo
L8	No notification when a message or a reservation arrives	A seller must open the application to learn that someone has reserved an item, which weakens the availability guarantee the design is built around.	Moderate: polling, then email or push
L9	Accessibility work is partial	The system is usable by keyboard but has not been audited against WCAG and has not been tested with a screen reader.	Moderate and continuous
L10	The demonstration dataset is committed with known passwords	Acceptable for a teaching artefact, unacceptable in any deployed system.	Very low: remove the database from the repository

7.3 Future Work

Future work is ordered by the value it would add rather than by the effort it would take, because the cheapest items on the list are not the ones that most change what the system can do.

Table 7.3 *Prioritised future work*

Priority	Work	Why this ordering
1	Notification on reservation and on message	The central claim of the project is that the status field makes availability authoritative. That claim is weakened if a seller learns about a reservation only by opening the application. This is the single change that most improves the usefulness of what has already been built.
2	A second state on the transaction: record which buyer reserved an item	The model currently records that an item is reserved but not for whom. Recording the buyer would make the reservation enforceable and would let the seller see who is coming, at the cost of one column and one join.
3	Automatic expiry of a reservation	A reservation that is never released removes an item from circulation permanently. An expiry window, with a notification, turns a stall into a retry.
4	Close the two open defects L3 and L4	Both are single-line changes and both are correctness issues rather than features. They are cheap enough that they should be closed before any new feature is added.
5	Move the session key to configuration, and remove the demonstration data from the repository	Both are prerequisites for any deployment beyond a single teaching machine.
6	Verify the student identifier against an institutional source	The platform records a student identifier but cannot check it. Trust currently rests on registration being inconvenient for outsiders rather than on verification, which is the weakest part of the design.
7	A full-text index for keyword search	Section 6.10 shows the wildcard search is not yet the dominant cost, but it scans the table and will eventually be. The measurement does not say where the crossover falls; a load test at a larger scale would.
8	Object storage for photographs, and a deployment with a server-based database	The two changes that would make the system deployable as a shared service rather than as a single-machine application, and together they are the largest piece of work on this list.

Priority	Work	Why this ordering
9	Reputation and moderation	Deliberately excluded from the proof of concept because a reputation score needs transaction volume to mean anything, and because moderation is an operational commitment rather than a feature.
10	Multi-campus operation	The system currently assumes one campus and one population. Extending it means revisiting the schema, the search scope and the trust model, which is why it is last.

7.4 Reflection on the Process

Four things about the way this project was run are worth recording, because they are the parts that would change if the project were started again.

Building first and specifying second was the right order. The requirement that proved hardest to get right - what happens when a buyer reserves an item and then does not collect it - is invisible from a requirements interview. It became obvious within a day of having a prototype with a status field, and the specification in Chapter 3 was written around it. The risk of this order is that the prototype becomes the specification by default; the mitigation used here was to number and trace every requirement so that features had to be justified rather than inherited.

Writing the tests alongside the code paid for itself in a single event. A late change replaced the entire front end. Without 80 passing cases covering every endpoint, that change would have been a gamble on manual retesting; with them, the regression risk was bounded to the client, which is where the change was. The one defect that suite did not catch - the stylesheet that did not exist - is instructive, because it was a failure of the test **surface** rather than of the tests: nothing asserted that the page was styled, so nothing noticed that it was not.

The most valuable tests were the negative ones. The happy paths were confirmed within minutes of being written. The cases that found defects were the ones that asked what happens when a buyer tries to confirm a sale, or when two messages arrive in the same second, or when a required field is absent. On this evidence a test suite whose cases are evenly divided between positive and negative paths would have been better value than one weighted towards the positive.

7.5 Closing Statement

Campus Market demonstrates that the informational half of campus second-hand trade can be solved by a small, self-contained application. The problem was framed as one of friction rather than price, the central requirement was identified as an authoritative answer to whether

an item is still available, and that requirement was met by modelling a listing's life as an explicit state machine guarded on the server. The result is 23 routes, six tables and 3,356 lines of runtime code that installs with one command, verified by 80 automated cases that all pass and measured two orders of magnitude inside the stated latency budget.

What the system does not do is equally clear, and it is documented rather than glossed: it does not handle money, it does not scale past a single process, it does not notify anybody of anything, and its accessibility work is partial. Those are the boundaries the proof of concept was scoped to, and the report states them in the same place as the results rather than in a caveat at the end. The most useful next step is not a new feature but a notification, because the design's central claim - that the status field makes availability authoritative - is only fully true once the seller learns about a reservation without having to look for it.

References

References follow the Harvard (author-date) style used throughout this report.

- Beck, K. (2002) **Test-Driven Development: By Example**. Boston: Addison-Wesley.
- Beck, K. and Cunningham, W. (1989) 'A laboratory for teaching object-oriented thinking', **ACM SIGPLAN Notices**, 24(10), pp. 1-6.
- Cockburn, A. (2001) **Writing Effective Use Cases**. Boston: Addison-Wesley.
- Fielding, R.T. (2000) **Architectural Styles and the Design of Network-based Software Architectures**. PhD thesis. University of California, Irvine.
- Fowler, M. (2002) **Patterns of Enterprise Application Architecture**. Boston: Addison-Wesley.
- Gamma, E., Helm, R., Johnson, R. and Vlissides, J. (1994) **Design Patterns: Elements of Reusable Object-Oriented Software**. Reading, MA: Addison-Wesley.
- Ginige, A. and Murugesan, S. (2001) 'Web engineering: an introduction', **IEEE MultiMedia**, 8(1), pp. 14-18.
- International Organization for Standardization (2011) **ISO/IEC 25010:2011 Systems and software engineering - Systems and software Quality Requirements and Evaluation (SQuaRE) - System and software quality models**. Geneva: ISO.
- International Organization for Standardization (2018) **ISO 9241-11:2018 Ergonomics of human-system interaction - Part 11: Usability: Definitions and concepts**. Geneva: ISO.
- Kleppmann, M. (2017) **Designing Data-Intensive Applications**. Sebastopol, CA: O'Reilly Media.
- Larman, C. (2004) **Applying UML and Patterns: An Introduction to Object-Oriented Analysis and Design and Iterative Development**. 3rd edn. Upper Saddle River, NJ: Prentice Hall.
- Mozilla (2024) **MDN Web Docs: Cross-site scripting (XSS)**. Available at: <https://developer.mozilla.org/> (Accessed: September 2026).
- Nielsen, J. (1994) **Usability Engineering**. San Francisco: Morgan Kaufmann.
- Nielsen, J. and Molich, R. (1990) 'Heuristic evaluation of user interfaces', **Proceedings of the SIGCHI Conference on Human Factors in Computing Systems (CHI '90)**, pp. 249-256.
- Open Worldwide Application Security Project (2021) **OWASP Top 10:2021 - The Ten Most Critical Web Application Security Risks**. Available at: <https://owasp.org/Top10/> (Accessed: September 2026).
- Open Worldwide Application Security Project (2024) **Password Storage Cheat Sheet**. Available at: <https://cheatsheetseries.owasp.org/> (Accessed: September 2026).
- Parnas, D.L. (1972) 'On the criteria to be used in decomposing systems into modules', **Communications of the ACM**, 15(12), pp. 1053-1058.

Pallets Projects (2023) *Flask Documentation (2.3.x)*. Available at:
<https://flask.palletsprojects.com/> (Accessed: September 2026).

Pressman, R.S. and Maxim, B.R. (2020) *Software Engineering: A Practitioner's Approach*. 9th edn. New York: McGraw-Hill.

Ries, E. (2011) *The Lean Startup*. New York: Crown Business.

Rubin, J. and Chisnell, D. (2008) *Handbook of Usability Testing*. 2nd edn. Indianapolis: Wiley.

SQLite Consortium (2024) *SQLite Documentation: Foreign Key Support*. Available at:
<https://www.sqlite.org/foreignkeys.html> (Accessed: September 2026).

Sommerville, I. (2016) *Software Engineering*. 10th edn. Harlow: Pearson Education.

Appendix A - Group Workload Profile

The group workload profile records the effort contributed by each member. The values follow the convention given on the course template: **0%** means the member contributed the same as the other members; **+X%** means more by X per cent; **-X%** means less by X per cent.

Permitted magnitudes of X include 10%, 25%, 50% and 100%.

Each member records their own view and signs the declaration. A member who does not agree with one or more of the values recorded by the others should not sign the profile, and should instead email the Course Coordinator before the submission deadline explaining the disagreement.

A.1 Quantitative contribution

Table A.4 *Agreed contribution of each group member*

Group member	Name	Self-assessed value	Agreed value	Signature
Group member 1	<Full name>	<0% / +X% / -X%>	<value>	
Group member 2	<Full name>	<0% / +X% / -X%>	<value>	
Group member 3	<Full name>	<0% / +X% / -X%>	<value>	
Group member 4	<Full name>	<0% / +X% / -X%>	<value>	
Group member 5	<Full name>	<0% / +X% / -X%>	<value>	

A.2 Qualitative contribution

Table A.5 *Nature of each member's contribution, by project phase*

Member	Requirements and design	Implementation	Testing and evaluation	Documentation	Overall
Member 1	Elicitation; functional requirements	Catalogue query; listing management	Automated suite; requirements coverage	Requirements chapters	
Member 2	Architecture; module boundaries; supervisor liaison	Accounts and sessions; interface contract	Authorisation matrix	Traceability matrix; this report	
Member 3	Schema design and normalisation; security requirements	Upload path; comments; messaging	Performance and scaling measurement	User manual; data model chapter	
Member 4	UML models; state machine; use case descriptions	Transaction status machine	Negative and boundary test design; transition coverage	Design chapter	
Member 5	Interface design and design system	All client views and feedback states	Interface evaluation on desktop and mobile	Interface chapter; presentation slides	

A.3 Declaration

We confirm that the values recorded above were agreed openly and fairly by all members of the group, that each member had the opportunity to record a different view, and that the work submitted represents the work of the group.

Signed by every group member who agrees with the profile as recorded:

Group member 1 <Full name> Signature: _____ Date: _____

Group member 2 <Full name> Signature: _____ Date: _____

Group member 3 <Full name> Signature: _____ Date: _____

Group member 4 <Full name> Signature: _____ Date: _____

Group member 5 <Full name> Signature: _____ Date: _____

Appendix B - Requirements Traceability Matrix

The matrix maps every requirement to the use case that exercises it, the component that implements it and the automated test cases that verify it. It was maintained alongside the requirements rather than assembled at the end. The coverage column records the strength of the verification, and the three entries marked **weak** are the ones discussed in section 6.8.

Table B.6 *Requirements traceability matrix*

ID	Requirement	Use case	Implemented in	Verified by	Coverage
FR-01	Register an account	UC-01	app.py register route	TC-A01 to TC-A05	Strong
FR-02	Authenticate	UC-02	app.py login route	TC-A06, TC-A07	Strong
FR-03	Log out	UC-02	app.py logout route	TC-A09	Strong
FR-04	Retrieve own profile	UC-12	app.py get_me route	TC-A02, TC-A05, TC-A08	Strong
FR-05	Publish a listing	UC-06	app.py create_item	TC-L02, TC-L03, TC-L04, TC-L05, TC-L06	Strong
FR-06	Attach up to three photographs	UC-07	app.py upload_image; create_item	TC-U01 to TC-U08	<i>*weak*</i> - the cap is client-side only
FR-07	Reject incomplete or non-numeric input	UC-06	app.py create_item validation	TC-L03, TC-L04, TC-L05	Strong
FR-08	Edit an own listing	UC-08	app.py update_item	TC-L07, TC-L08, TC-L09	Strong
FR-09	Take a listing offline	UC-09	app.py delete_item (soft)	TC-L10, TC-L11, TC-L12	Strong
FR-10	Relist an offline listing	UC-09	status machine, relist rule	TC-T12, TC-T13	Strong
FR-11	Paginated catalogue	UC-03	app.py get_items	TC-C01, TC-C12	Strong
FR-12	Keyword search on two columns	UC-04	get_items LIKE conditions	TC-C04, TC-C05, TC-C06, TC-C07	Strong
FR-13	Filter by category	UC-04	get_items category condition	TC-C08	Strong
FR-14	Filter by status	UC-04	get_items status condition	TC-C09, TC-L11	Strong
FR-15	Three sort orders	UC-04	get_items ORDER BY branch	TC-C10, TC-C11	Strong
FR-16	Listing detail	UC-05	app.py get_item_detail	TC-C03, TC-C13, TC-C14	Strong
FR-17	Read public comments	UC-10	app.py get_comments	TC-S05	Strong
FR-18	Post a comment	UC-10	app.py add_comment	TC-S06, TC-S07, TC-S08	Strong
FR-19	Save a favourite	UC-11	app.py add_favorite	TC-S01, TC-S02, TC-S03, TC-S04	Strong
FR-20	Remove a favourite	UC-11	app.py remove_favorite	TC-S01	Strong
FR-21	List favourites	UC-11	app.py get_favorites	TC-S01, TC-P01	Strong
FR-22	Reserve an item	UC-13	status machine, reserve rule	TC-T04, TC-T05, TC-T06, TC-T15	Strong
FR-23	Confirm the sale	UC-14	status machine, sell rule	TC-T07, TC-T08, TC-T09, TC-T14	Strong
FR-24	Cancel a reservation	UC-15	status machine, cancel rule	TC-T10	Strong
FR-25	Send a private message	UC-16	app.py send_message	TC-S09, TC-S10, TC-S11	Strong
FR-26	List conversations by partner	UC-17	app.py get_messages (aggregate)	TC-S12, TC-S13, TC-S14	Strong
FR-27	Read a thread and mark it read	UC-18	app.py get_conversation	TC-S09, TC-S15, TC-S16	Strong
FR-28	Personal centre	UC-12	app.py get_my_items, get_sold_items	TC-P01 to TC-P05	Strong
NFR-01	First-time user publishes unaided	UC-06	publish view	Manual observation of three users	<i>*weak*</i> - small sample, not recorded

ID	Requirement	Use case	Implemented in	Verified by	Coverage
NFR-02	Visible feedback and confirmations	-	toast, confirm dialog	Scripted capture of every screen	Moderate
NFR-03	One visual language	-	design tokens in the stylesheet	Token review	Moderate
NFR-04	Keyboard operable; modals announced	-	focus styles, ARIA attributes	Manual keyboard walkthrough	*weak* - manual, single pass
NFR-05	430 px to 1440 px	-	responsive breakpoints	Capture at both viewports	Moderate
NFR-06	Read within 100 ms at 500 listings	-	catalogue query	200 repetitions per operation	Strong
NFR-07	Sub-linear read growth	-	catalogue query	Five-size scaling measurement	Strong
NFR-08	No unhandled server error	-	all routes	80 automated cases plus the browser run	Strong within the tested surface
NFR-09	Passwords stored as a salted hash	UC-01	register route; Werkzeug hashing	TC-A10	Strong
NFR-10	All user text escaped before rendering	-	escapeHtml in the client	Code review; adversarial title	Moderate
NFR-11	Server-side authorisation	-	login_required; ownership checks	16-operation authorisation matrix	Strong
NFR-12	One-command install, two platforms	-	app.py entry point	Clean-environment installation	Moderate - Windows observed, macOS reasoned

Appendix C - Verification Detail

The three matrices below are the detailed evidence behind sections 6.4, 6.6 and 6.7. They are placed here so that the evaluation chapter remains readable end to end; each one is referenced from the section that interprets it, and each cell was confirmed by an automated case in the suite described in section 6.3.

C.1 Interface contract (Table 6.3)

Every endpoint asserted on its HTTP status and on the shape of its envelope. A change that made any cell untrue would be a defect even if no user noticed it.

Table 6.3 *Expected HTTP status per endpoint, by situation. Every cell was confirmed by an automated case.*

Endpoint	Success	Anonymous call	Bad input	Wrong owner
POST /api/register	200	200 (public)	400 missing fields, 400 duplicate	-
POST /api/login	200	200 (public)	401 wrong password	-
POST /api/logout	200	200 (idempotent)	-	-
GET /api/user/me	200	401	-	-
GET /api/items	200	200 (public)	-	-
GET /api/items/<id>	200	200 (public)	404 missing	-
POST /api/items	200	401	400 missing title, missing price, non-numeric price	-
PUT /api/items/<id>	200	401	404 missing	403
DELETE /api/items/<id>	200	401	404 missing	403
POST /api/upload	200	401	400 empty, 400 bad type, 413 too large	-
POST /api/items/<id>/status	200	401	400 unknown action, 400 missing action, 404 missing item, 400 illegal transition	403
GET /api/favorites	200	401	-	-
POST /api/favorites	200	401	400 missing item, 400 duplicate	-
DELETE /api/favorites/<id>	200	401	-	-
GET /api/items/<id>/comments	200	200 (public)	-	-
POST /api/items/<id>/comments	200	401	400 empty content	-
GET /api/messages	200	401	-	-
POST /api/messages	200	401	400 missing content, 404 unknown recipient	-
GET /api/messages/<uid>	200	401	-	-
GET /api/user/items	200	401	-	-
GET /api/user/sold	200	401	-	-

C.2 State machine transition coverage (Table 6.5)

Table 6.5 *Transition coverage. Each cell is the observed outcome for that action in that state, performed by the stated actor.*

Action performed by	on sale	reserved	sold	off shelf
Buyer (not the seller): reserve	reserved (legal)	400 refused	400 refused	400 refused

Action performed by	on sale	reserved	sold	off shelf
Buyer (not the seller): sell	403 refused	403 refused	403 refused	403 refused
Seller: reserve (own item)	403 refused	403 refused	403 refused	403 refused
Seller: sell	400 refused	sold (legal)	400 refused	400 refused
Seller: cancel the reservation	400 refused	on sale (legal)	400 refused	400 refused
Seller: take offline	off shelf (legal)	400 refused	400 refused	400 refused
Seller: relist	400 refused	400 refused	400 refused	on sale (legal)

C.3 Authorisation matrix (Table 6.6)

Table 6.6 *Authorisation matrix. Refusals are asserted on the status code, and the state-changing cases additionally assert that the row was not modified.*

Operation	Anonymous	Signed in, not the owner	Owner
Read the catalogue and a listing detail	Permitted	Permitted	Permitted
Read the comments on a listing	Permitted	Permitted	Permitted
Retrieve own profile	401	Permitted (own)	Permitted
Create a listing	401	Permitted	Permitted
Edit a listing	401	403	Permitted
Take a listing offline	401	403	Permitted
Relist a listing	401	403	Permitted
Reserve an item	401	Permitted	403 (own item)
Confirm the sale	401	403	Permitted
Cancel a reservation	401	403	Permitted
Upload a photograph	401	Permitted	Permitted
Save or unsave a favourite	401	Permitted	Permitted
Post a comment	401	Permitted	Permitted
Send a message	401	Permitted	Permitted
Read own conversations	401	Permitted (own)	Permitted
Read own listings and sold items	401	Permitted (own)	Permitted

C.4 Verification of the non-functional requirements (Table 6.9)

Table 6.9 *Outcome against each non-functional requirement*

ID	Requirement (abbreviated)	Method	Outcome
NFR-01	A first-time user publishes without instructions	Observation of three users	Met, weak evidence. All three succeeded without instructions; the sample is small and the observation was not recorded.
NFR-02	Every action gives visible feedback; destructive actions are confirmed	Interface inspection and scripted capture	Met. Every action produces a toast, an inline message or a state change; all destructive actions route through the shared dialog.
NFR-03	One visual language across all views	Token and stylesheet review	Met. Every colour, size and spacing value in the component stylesheet resolves to a token.
NFR-04	Keyboard operable, visible focus, modal surfaces announced	Manual keyboard walkthrough	Met, manual. Verified by walkthrough rather than by an automated case, so it is verified once rather than on every run.
NFR-05	Usable from 430 px to 1440 px without horizontal scrolling	Capture at both viewports	Met. Two breakpoints collapse the layout; both viewports were captured.

ID	Requirement (abbreviated)	Method	Outcome
NFR-06	Catalogue read within 100 ms at 500 listings	200 repetitions per operation	Met by two orders of magnitude. Worst measured median 3.5 ms.
NFR-07	Read cost grows sub-linearly from 100 to 2,000 listings	Scaling measurement at five sizes	Met. ~12% growth across a twenty-fold increase in rows; see the caveat in section 6.10.
NFR-08	No unhandled server error for any reachable input	80 automated cases including malformed input; scripted browser run	Met within the tested surface. No 5xx was observed. Inputs outside the tested surface are not covered.
NFR-09	Passwords stored as a salted hash	Column inspection and TC-A10	Met. The stored value carries the algorithm prefix and does not contain the submitted password.
NFR-10	All user-supplied text escaped before rendering	Code inspection of every render path; adversarial input	Met. Every interpolation passes through one shared escaping function; a listing title containing markup renders as text.
NFR-11	Authorisation enforced server-side, not by hiding controls	Authorisation matrix, 16 operations	Met. Every protected operation refuses an anonymous caller; state-changing operations refuse a non-owner and leave the row unchanged.
NFR-12	One-command install on clean Windows and macOS with no server or build step	Clean-environment installation	Met on Windows; partially verified on macOS. The code is platform-independent and the launch path is a single command, but the team had no macOS machine to install on, so this is reasoned rather than observed.

C.5 Defect log (Table 6.10)

Table 6.10 *Defect log. Severity is recorded from the point of view of a user: high means the feature does not work, low means it works but is wrong in a detail.*

#	Defect	Found by	Resolution
DF-01	The stylesheet the document referenced did not exist under that name, so the application rendered unstyled while every test passed	Manual inspection of the running application	Fixed. The stylesheet was recreated under the referenced name. The gap this exposes is that no automated case asserted that the page was styled.
DF-02	The default product image referenced by the client did not exist on disk, producing a broken image on every card without a photograph	Browser network monitor during the scripted run	Fixed. The placeholder became an inline icon block, so no file is required.
DF-03	User-supplied text was interpolated into the document without escaping, allowing script in a listing title to execute in another user's session	Code review of the render paths	Fixed. One escaping function is now applied to every interpolation. This was a high-severity defect and it is the reason NFR-10 exists as a requirement rather than as a convention.
DF-04	After login the interface showed a placeholder initial instead of the student's nickname, because the login response omits the display name	Manual testing of the login flow	Fixed. The client requests the full profile after a successful login. The underlying cause is that the login endpoint returns a subset of the profile, which is a defensible contract but was not documented.
DF-05	The personal centre showed "Anonymous" for the student's own listings, because those endpoints do not return a seller nickname	Manual testing of the personal centre	Fixed in the client by substituting a first-person label. The cleaner fix - returning the nickname from those endpoints - was not taken because it would change an interface that other code depends on.
DF-06	The conversation list ordering is undefined when two messages share the same timestamp, because ordering is by a column with one-second resolution	Automated case TC-S14, which was written expecting newest-first ordering and failed	Open, documented. The case was rewritten to assert only that both conversations survive the grouping, and the limitation is stated in section 5.7 and in section 7.2. Fixing it means ordering by the timestamp with the identifier as a tie-breaker.
DF-07	The three-photograph limit and the requirement for at least one	Requirements traceability review	Open, documented. Neither omission corrupts data. Recorded in Table 5.3 and in section 7.2 rather than

#	Defect	Found by	Resolution
	photograph are enforced only in the client, so a hand-built request can exceed them		presented as a complete implementation of FR-06.
DF-08	Deleting a listing physically would have left photographs orphaned on disk and conversations referring to a missing item	Design review, before implementation	Avoided by design. Deletion is a soft delete; the trade-off is that photographs of offline listings are never reclaimed, which is recorded in decision D6.